



# Virtual Memory

Arkaprava Basu

Dept. of Computer Science and Automation  
Indian Institute of Science  
[www.csa.iisc.ac.in/~arkapravab/](http://www.csa.iisc.ac.in/~arkapravab/)




# Acknowledgements

- Some of the slides in the deck were provided by Profs. Luis Ceze (Washington), Nima Horanmand (Stony Brook), Mark Hill, David Wood, Karu Sankaralingam (Wisconsin), Abhishek Bhattacharjee (Yale).
- Development of this course is partially supported by funding from Western Digital corporations.

# Memory is virtual !

- Application software sees **virtual** address
  - ▶ `int * p = malloc(64);`
    - ← Virtual address
  - ▶ Load, store instructions carries virtual addresses

# Memory is virtual !

- Application software sees **virtual** address
  - ▶ `int * p = malloc(64);`  
 Virtual address
  - ▶ Load, store instructions carries virtual addresses
- Hardware uses (real) **physical** address
  - ▶ e.g., to find data, lookup caches etc.
- When an application executes (a.k.a process) virtual addresses are translated to physical addresses, at runtime

# Bird's view of virtual memory



Physical Memory (e.g., DRAM)

# Bird's view of virtual memory

Process 1

The diagram illustrates the 'Bird's view of virtual memory'. It consists of two main rectangular boxes. The top box, labeled 'Process 1', is outlined in green and is empty. The bottom box, labeled 'Physical Memory (e.g., DRAM)', is outlined in blue and is also empty. These boxes represent the virtual address space of a process and the physical memory it maps to, respectively.

Physical Memory (e.g., DRAM)

# Bird's view of virtual memory

Process 1

```
// Program expects (*x)  
// to always be at  
// address 0x1000  
int *x = 0x1000;
```

Physical Memory (e.g., DRAM)

# Bird's view of virtual memory

Process 1

Process 2

```
// Program expects (*x)  
// to always be at  
// address 0x1000  
int *x = 0x1000;
```

Physical Memory (e.g., DRAM)



# Bird's view of virtual memory

Process 1

Process 2


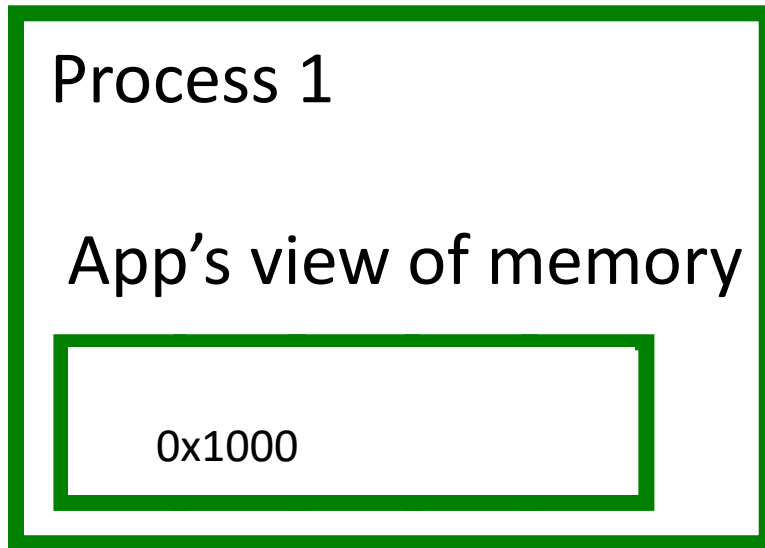
Only one physical  
address 0x1000!!

```
// Program expects (*x)  
// to always be at  
// address 0x1000  
int *x = 0x1000;
```

0x1000

Physical Memory (e.g., DRAM)

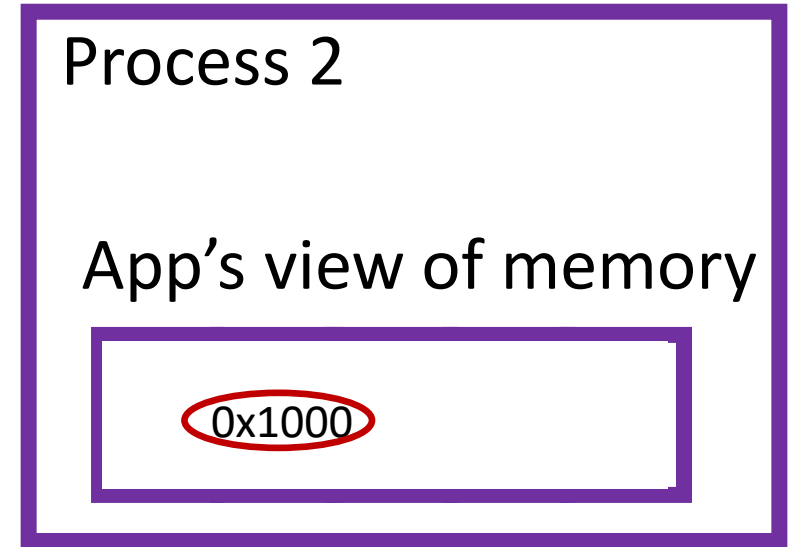

# Bird's view of virtual memory



0x1000 Physical Memory (e.g., DRAM)

# Bird's view of virtual memory

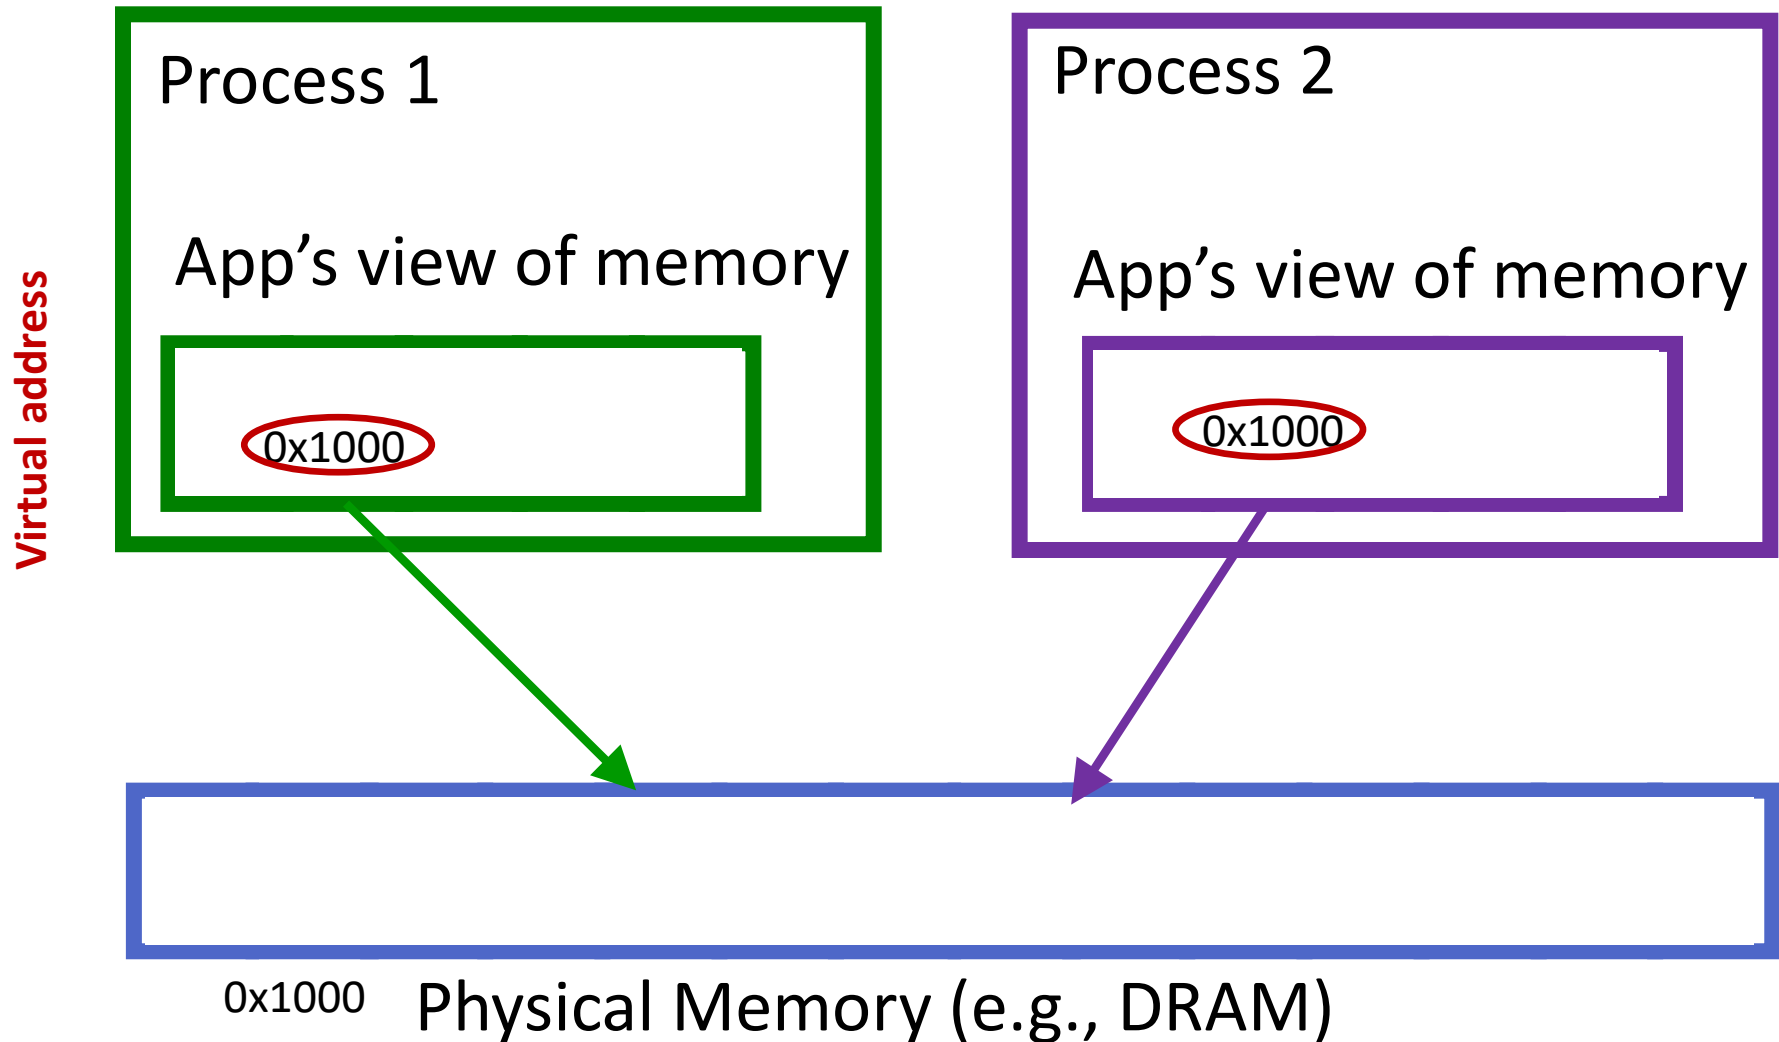
Virtual address



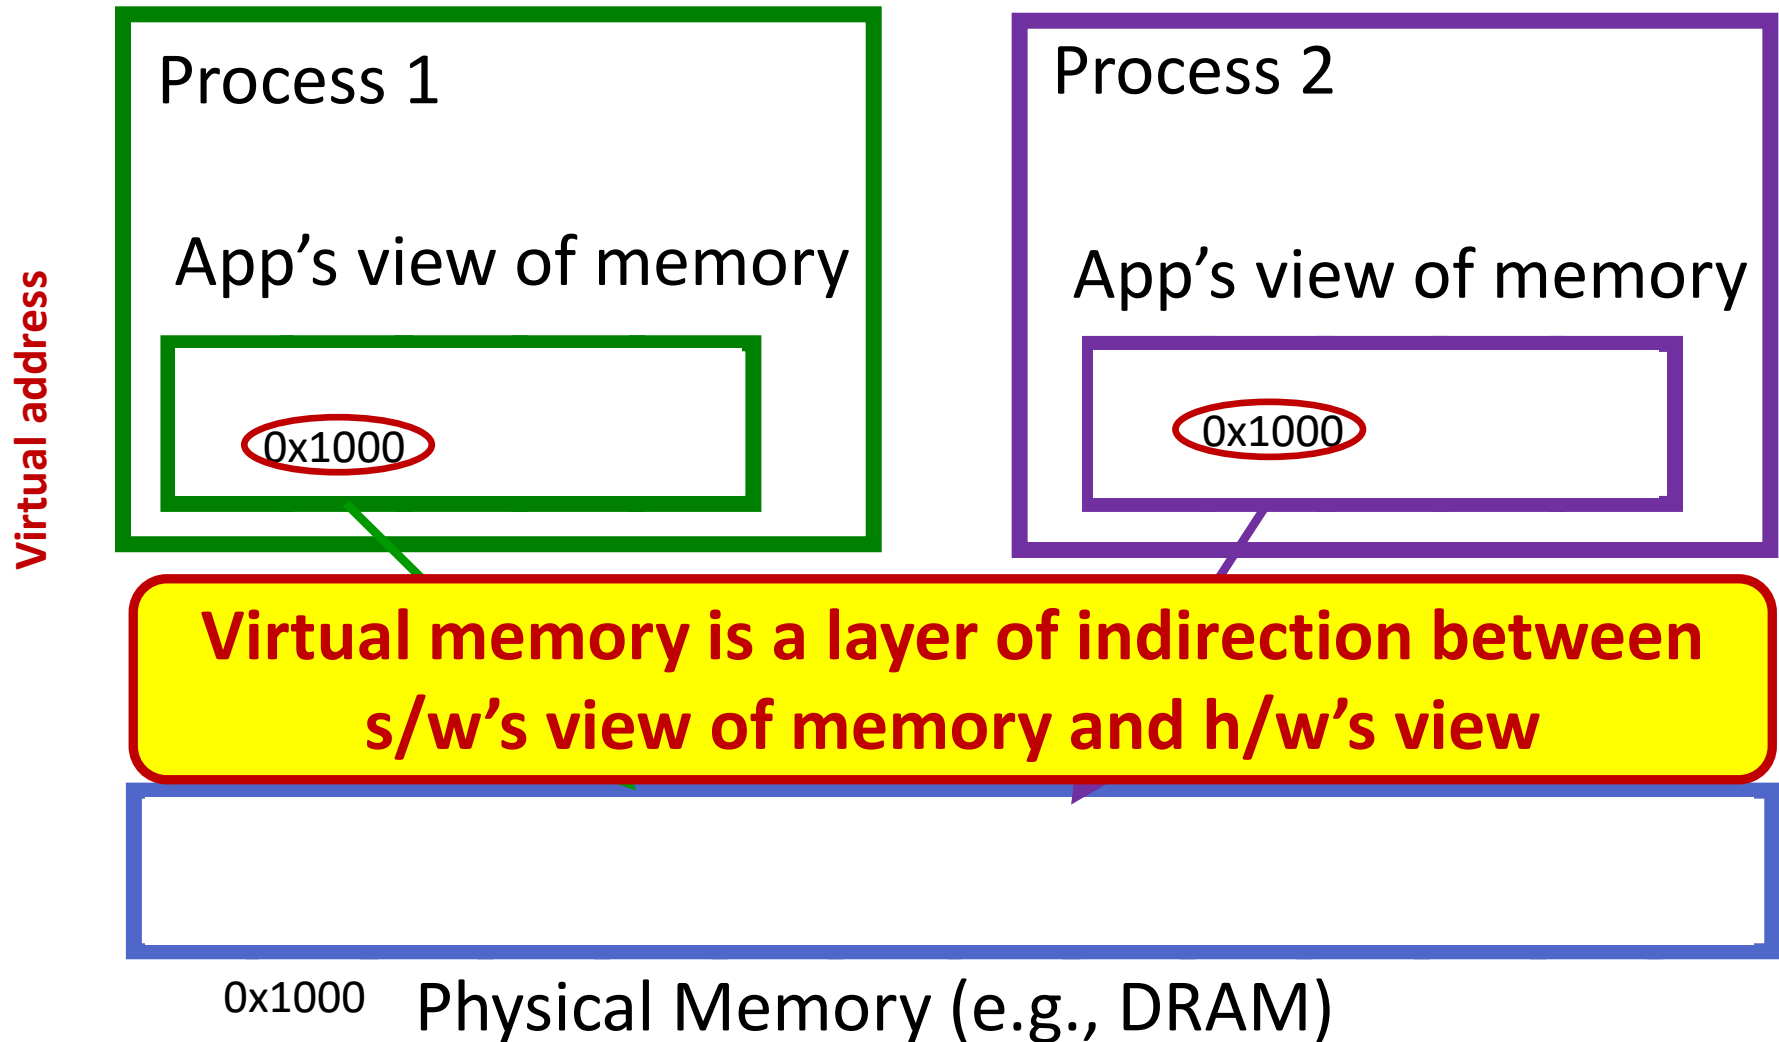
0x1000

Physical Memory (e.g., DRAM)

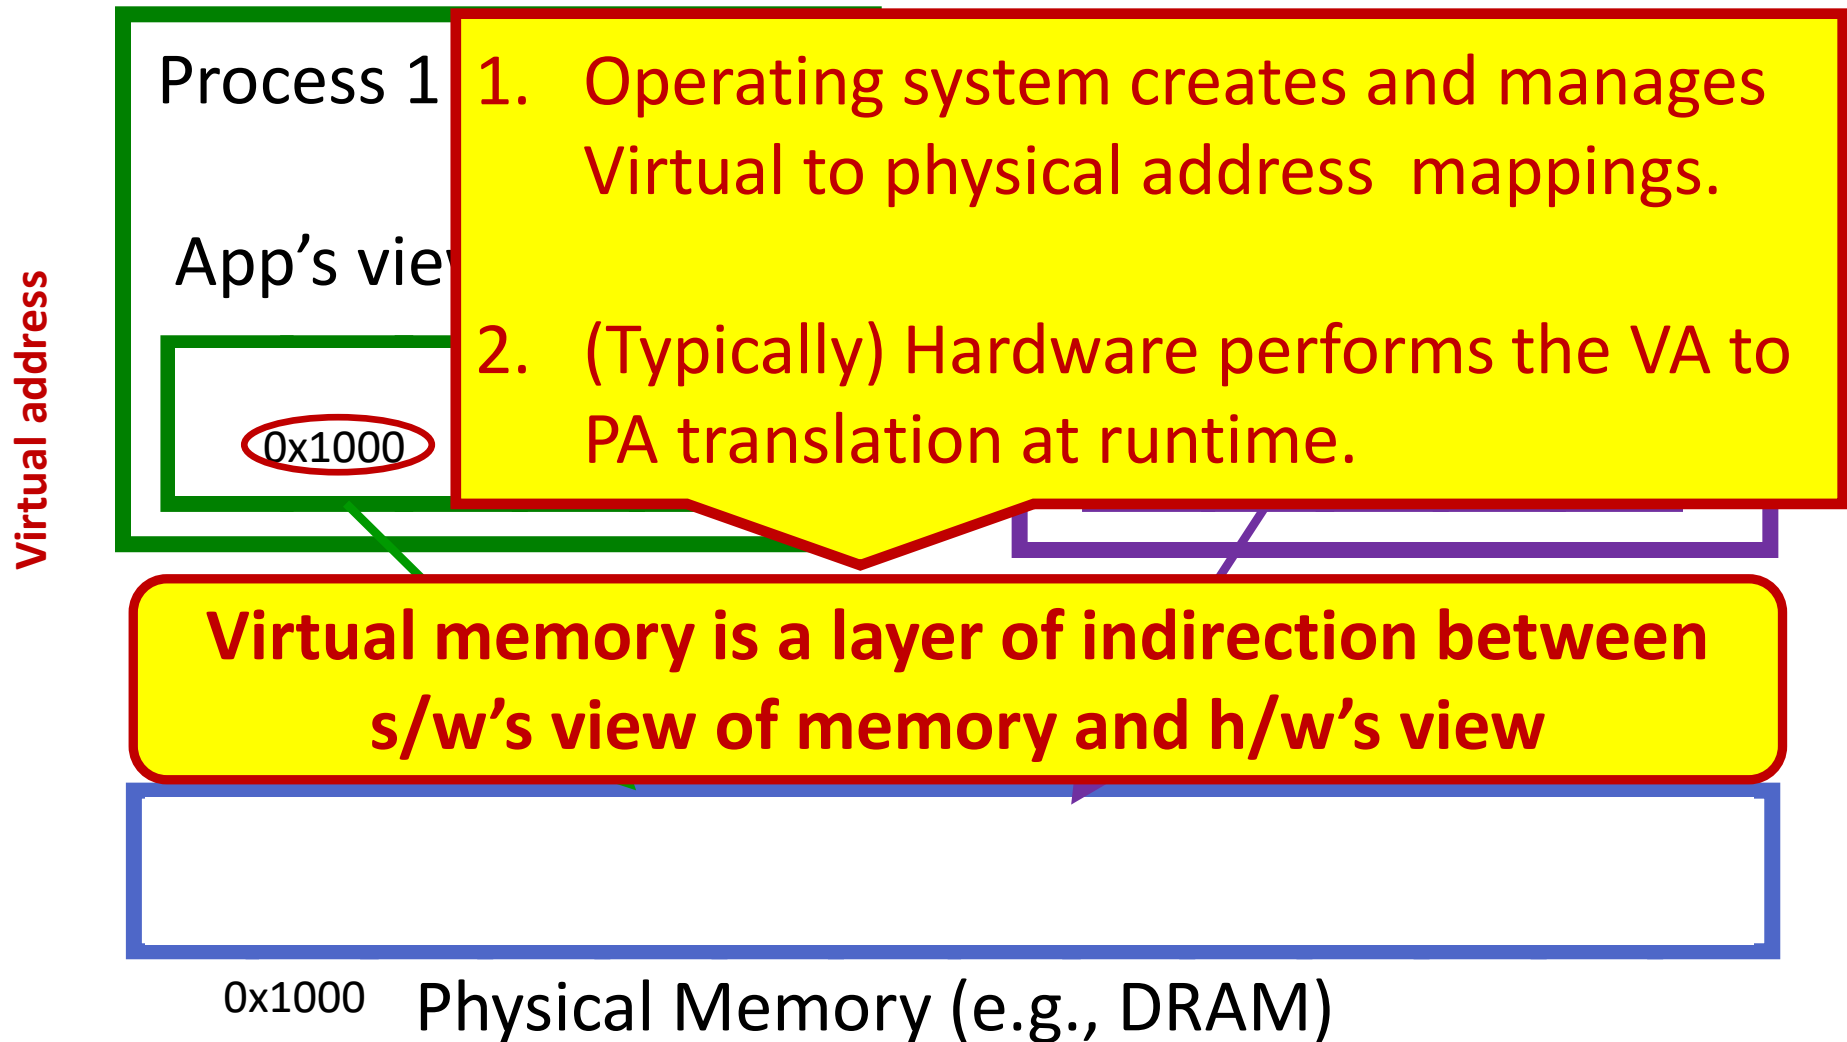
# Bird's view of virtual memory




# Bird's view of virtual memory




# Roles of software and hardware



# Virtual memory terminology

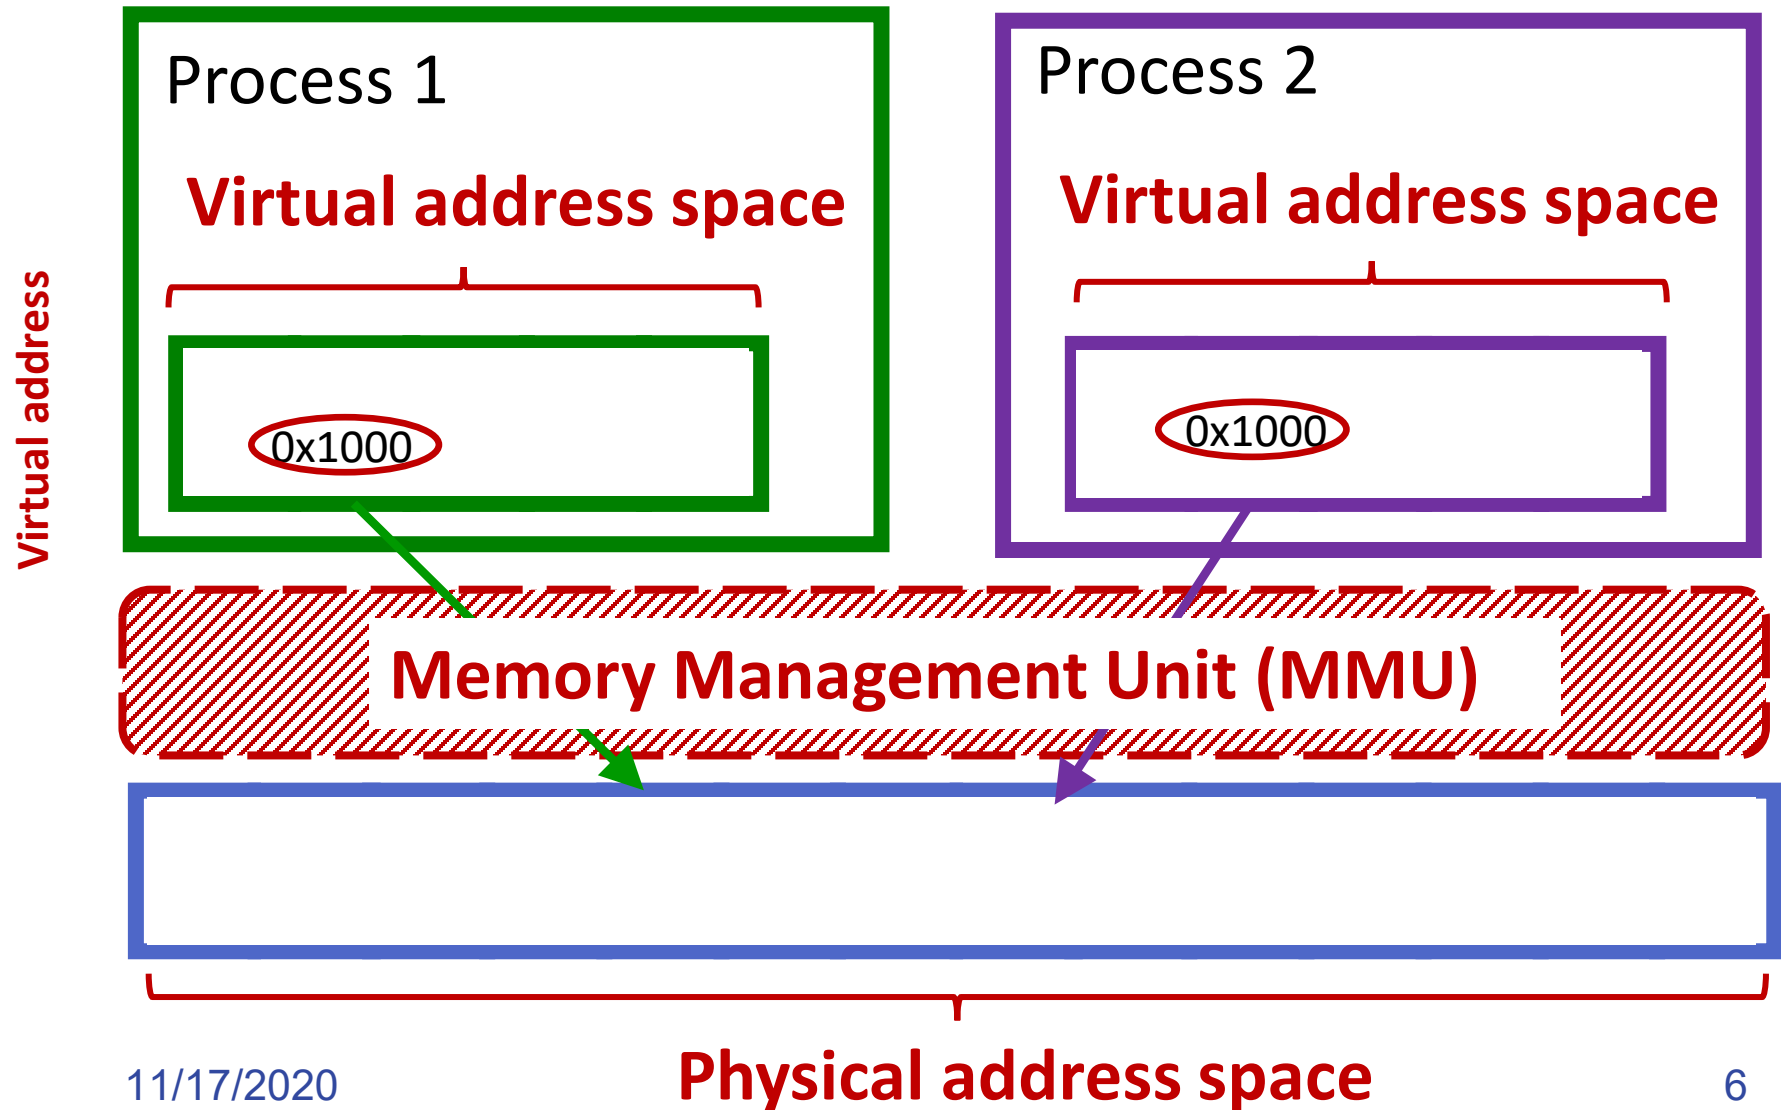


# Virtual memory terminology





# Virtual memory terminology





# Why virtual memory?



# Why virtual memory?

- **Relocation:** Ease of programming by providing application contiguous view of memory
  - ▶ But actual physical memory can be scattered



# Why virtual memory?

- **Relocation:** Ease of programming by providing application contiguous view of memory
  - ▶ But actual physical memory can be scattered
- **Resource management:** Application writer can assume there is enough memory
  - ▶ But, system can manage physical memory across concurrent processes
- **Isolation:** Bug in one process should not corrupt memory of another under multi-programming
  - ▶ Provide each process with separate virtual address space
- **Protection:** Enforce rules on what memory a process can or cannot access



# Why virtual memory?



# Why virtual memory?

- **Sharing/communication:** Inter-process communication and sharing

# Why virtual memory?

- **Sharing/communication:** Inter-process communication and sharing
  - ▶ Same physical address can be mapped onto two process's virtual address space with different virtual address



# Why virtual memory?

- **Sharing/communication:** Inter-process communication and sharing
  - ▶ Same physical address can be mapped onto two process's virtual address space with different virtual address
    - Could be used for inter-process-communication





# Why virtual memory?

- **Sharing/communication:** Inter-process communication and sharing
  - ▶ Same physical address can be mapped onto two process's virtual address space with different virtual address
    - Could be used for inter-process-communication
    - Very useful for code/library sharing



# Why virtual memory?

- **Sharing/communication:** Inter-process communication and sharing
  - ▶ Same physical address can be mapped onto two process's virtual address space with different virtual address
    - Could be used for inter-process-communication
    - Very useful for code/library sharing
- **Illusion of larger memory:** Allows memory overcommit by providing illusion of memory larger than actually available



# Why virtual memory?

- **Sharing/communication:** Inter-process communication and sharing
  - ▶ Same physical address can be mapped onto two process's virtual address space with different virtual address
    - Could be used for inter-process-communication
    - Very useful for code/library sharing
- **Illusion of larger memory:** Allows memory overcommit by providing illusion of memory larger than actually available
  - ▶ Can write program without worrying about amount of free memory in a system where the program will run

# Agenda

What is virtual memory?

Hardware implementations of virtual memory

Software management of virtual memory

Research opportunity in virtual memory

# Agenda

What is virtual memory?

Hardware implementations of virtual memory


Software management of virtual memory

Research opportunity in virtual memory

# Implementing virtual memory

- Many ways to implement virtual memory


- ▶ Base and bound register
- ▶ Segmentation
- ▶ Paging



# Base and bound registers

Process 1

Virtual address space



Process 2


Virtual address space



# Base and bound registers

Process 1

Virtual address space




Process 2

Virtual address space

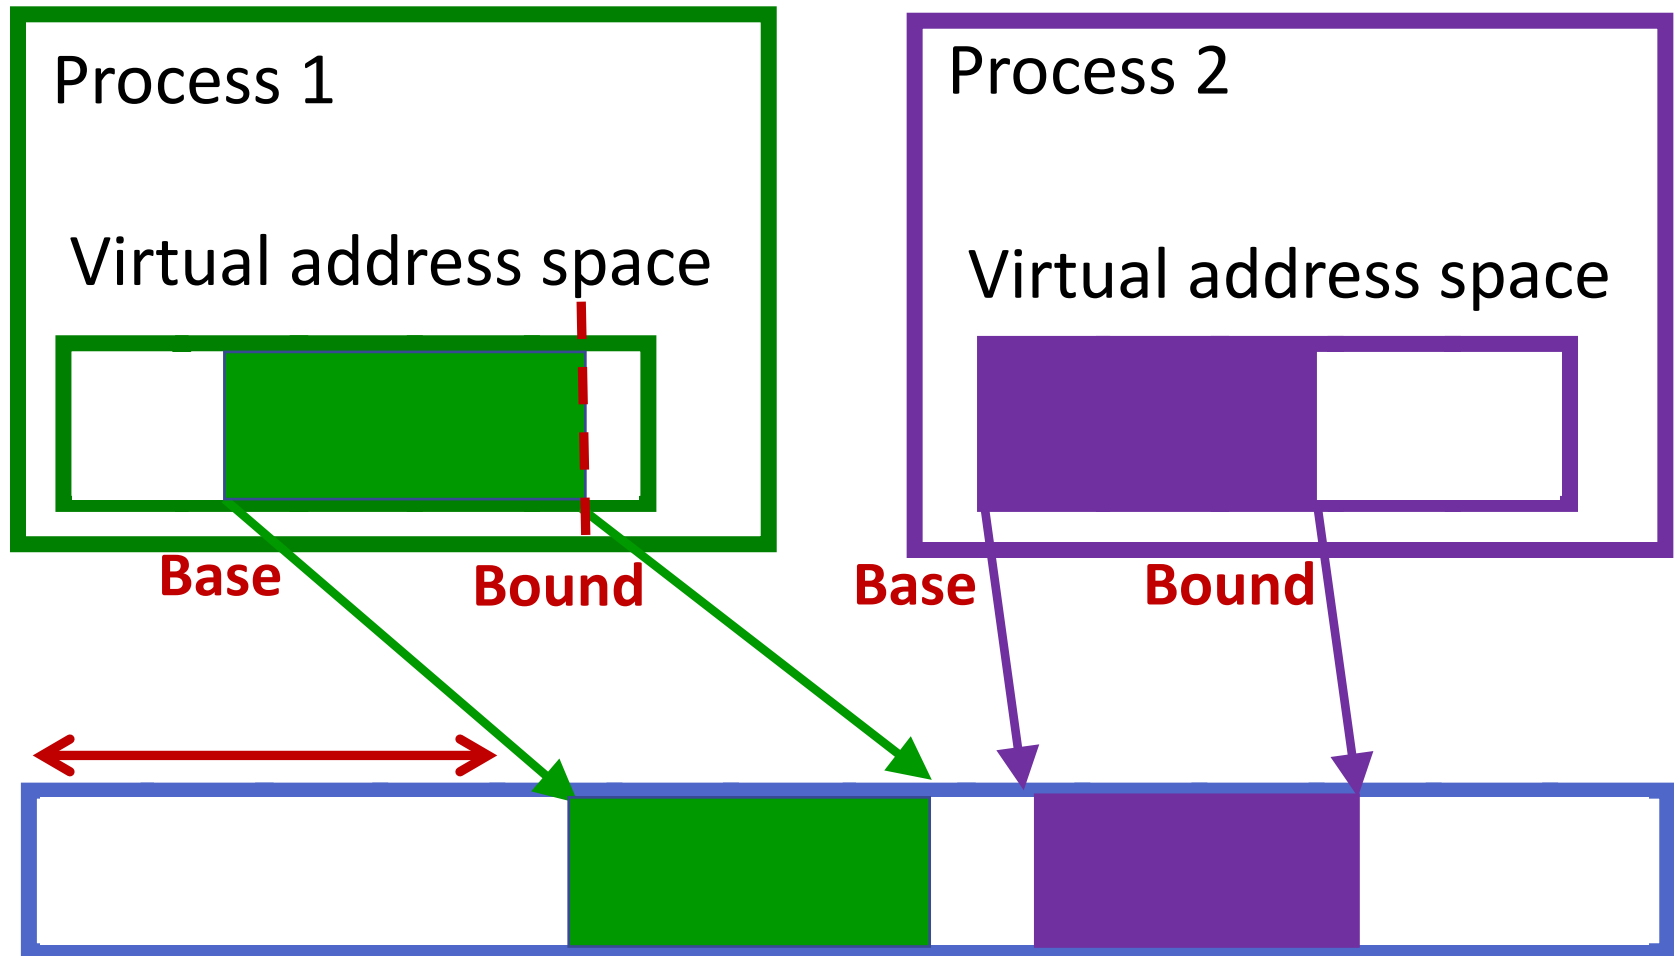




# Base and bound registers



# Base and bound registers



# Base and bound registers

- How it works?
  - ▶ Each CPU core has base and bound registers
  - ▶ Each process has own *values* for base and bound
  - ▶ Base and bound values are assigned by OS
  - ▶ On a context switch, a process's base and bound registers are saved/restored by the OS

# Base and bound registers

- How it works?
  - ▶ Each CPU core has base and bound registers
  - ▶ Each process has own *values* for base and bound
  - ▶ Base and bound values are assigned by OS
  - ▶ On a context switch, a process's base and bound registers are saved/restored by the OS
- Example: Cray-1 (very old system)
- Advantage
  - ▶ Simple


# Base and bound registers

- How it works?
  - ▶ Each CPU core has base and bound registers
  - ▶ Each process has own *values* for base and bound
  - ▶ Base and bound values are assigned by OS
  - ▶ On a context switch, a process's base and bound registers are saved/restored by the OS
- Example: Cray-1 (very old system)
- Advantage
  - ▶ Simple
- Disadvantage
  - ▶ Needs contiguous physical memory
  - ▶ Deallocation of memory is possible only on the edge
  - ▶ Cannot allocate more memory if no free memory at the edge

# Segmentation

Process 1

Virtual address space




Process 2


Virtual address space



# Segmentation

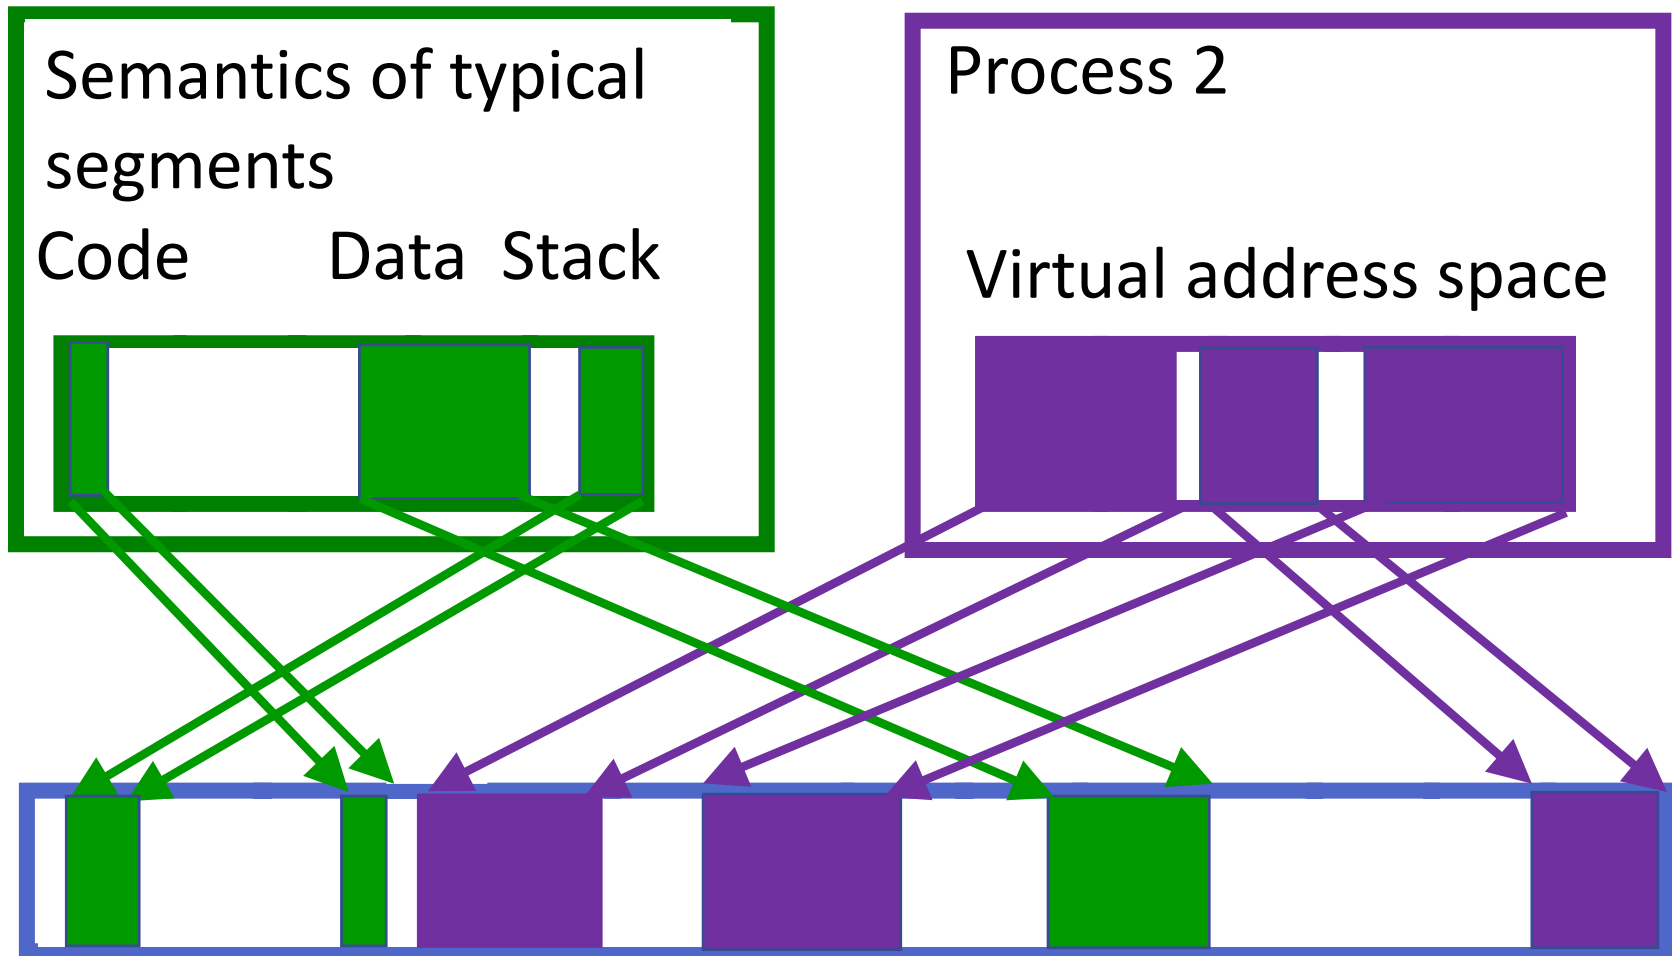


# Segmentation





# Segmentation



# Segmentation: How it works?

**Virtual address**

<b>CS</b>	<b>128</b>
-----------	------------

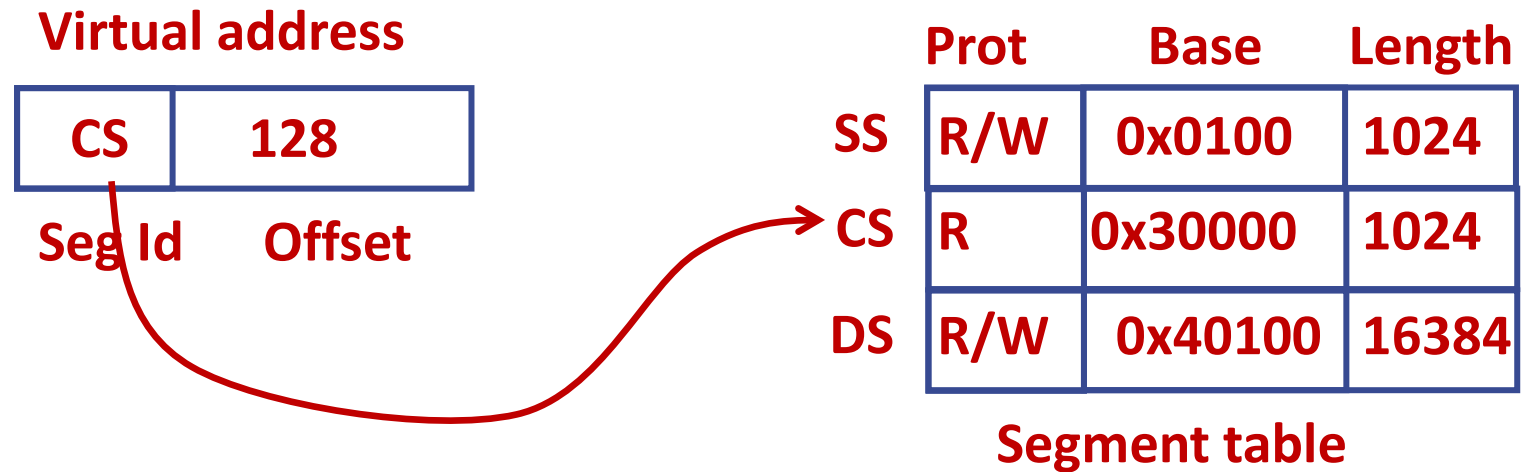
**Seg Id      Offset**

	<b>Prot</b>	<b>Base</b>	<b>Length</b>
<b>SS</b>	<b>R/W</b>	<b>0x0100</b>	<b>1024</b>
<b>CS</b>	<b>R</b>	<b>0x30000</b>	<b>1024</b>
<b>DS</b>	<b>R/W</b>	<b>0x40100</b>	<b>16384</b>

**Segment table**



# Segmentation: How it works?



# Segmentation: How it works?

Virtual address

CS	128
----	-----

Seg Id      Offset

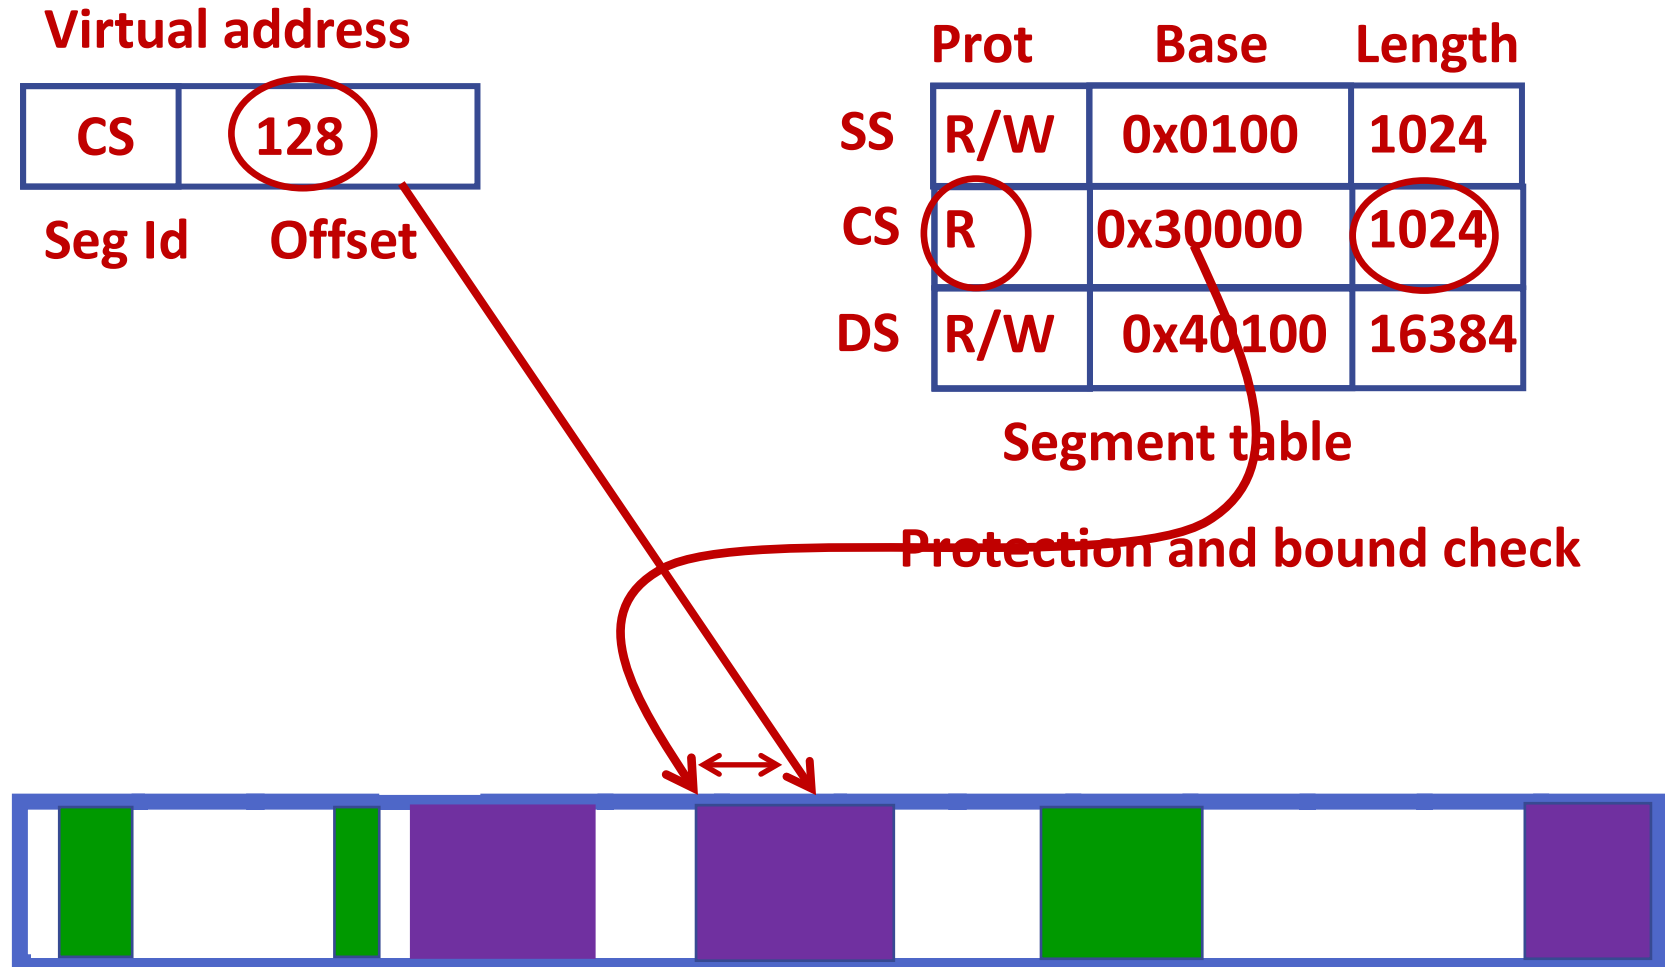
	Prot	Base	Length
SS	R/W	0x0100	1024
CS	R	0x30000	1024
DS	R/W	0x40100	16384

Segment table

Protection and bound check



# Segmentation: How it works?





# Segmentation: How it works?

## ■ How to select segmentation ID?

### ▶ In most cases, the compiler can infer


- For example, x86-32 bit has CS, SS, DS, ES, FS, GS

	
Compiler selected	Programmer selected


## ■ Advantages segmentation:

- ▶ Multiple memory regions with different protections
- ▶ No need to have all segments in memory all the time
- ▶ Ability to share segments across processes

# Disadvantages of Segmentation

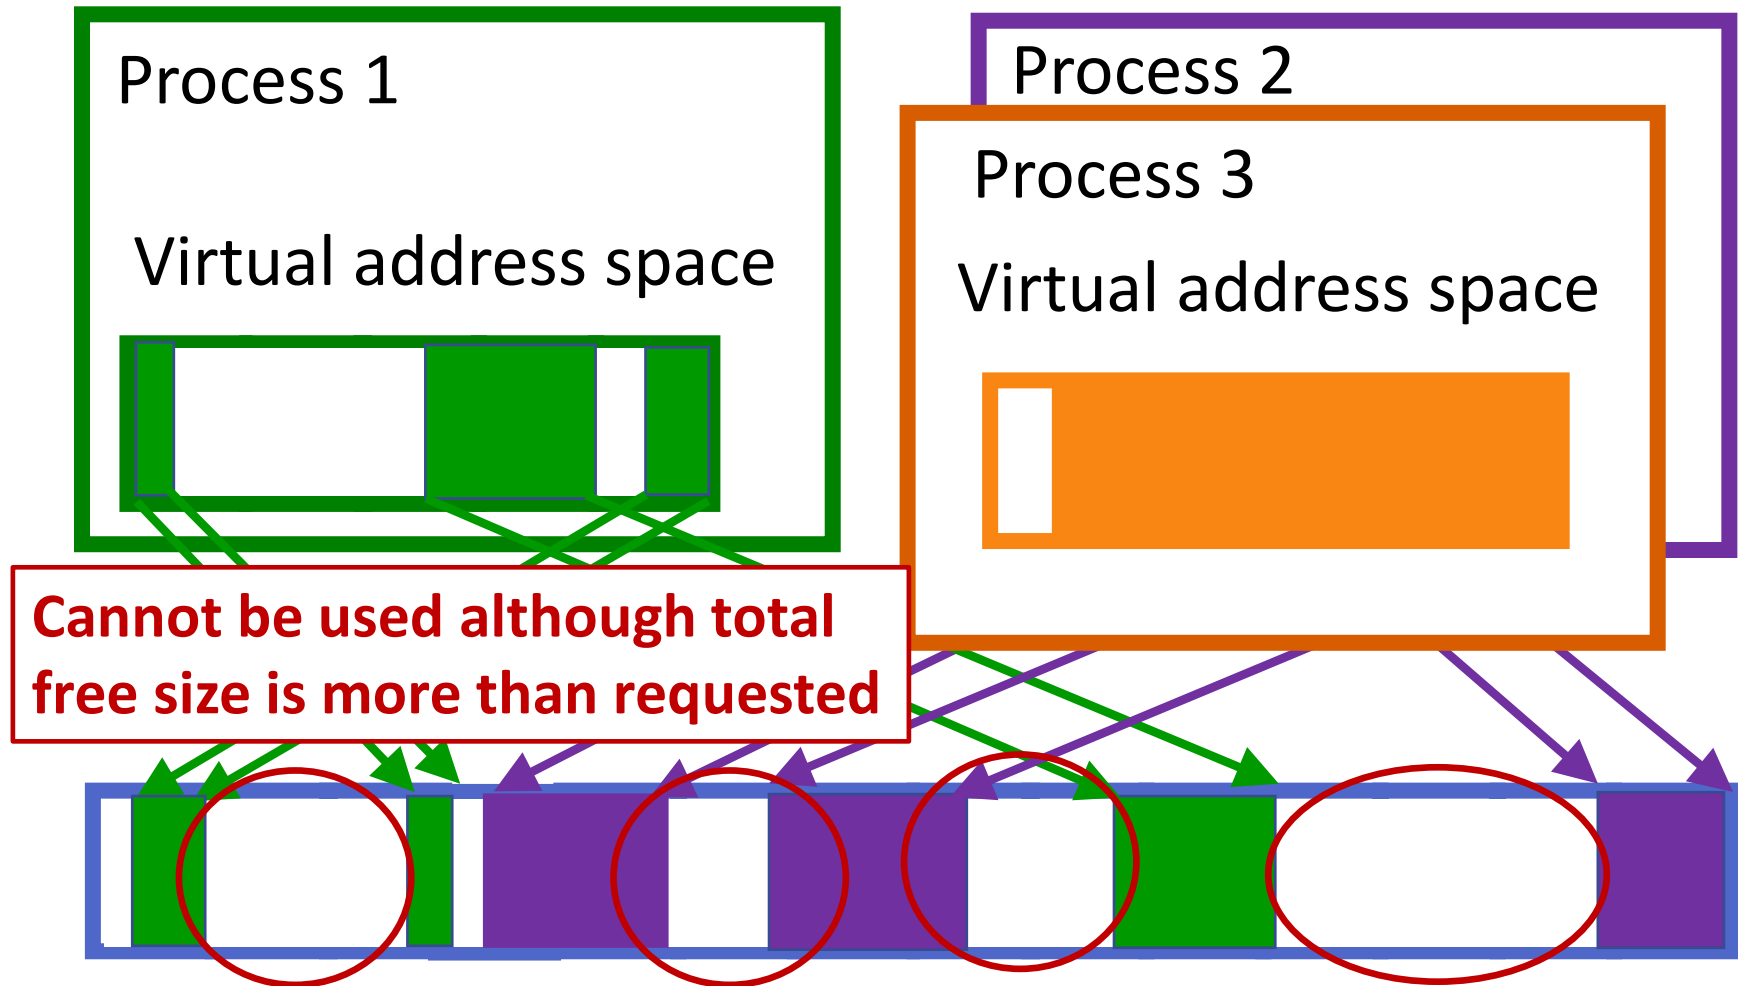


# Disadvantages of Segmentation





# Disadvantages of Segmentation





# Disadvantages of Segmentation

- **External fragmentation:** Inability to use free memory as it is non-contiguous (fragmentation)
  - ▶ Allocating different sized segments leaves free memory fragmented
- A segment of size  $n$  bytes requires  $n$  bytes long contiguous physical memory
- Not completely transparent to applications




# Paging

- **Idea:** Allocate/de-allocate memory in **same size** chunks (pages)
  - ▶ No external fragmentation (Why?)
  - ▶ Page sizes are typically small, e.g., 4KB

# Paging


Process 1

Virtual address space




Process 2

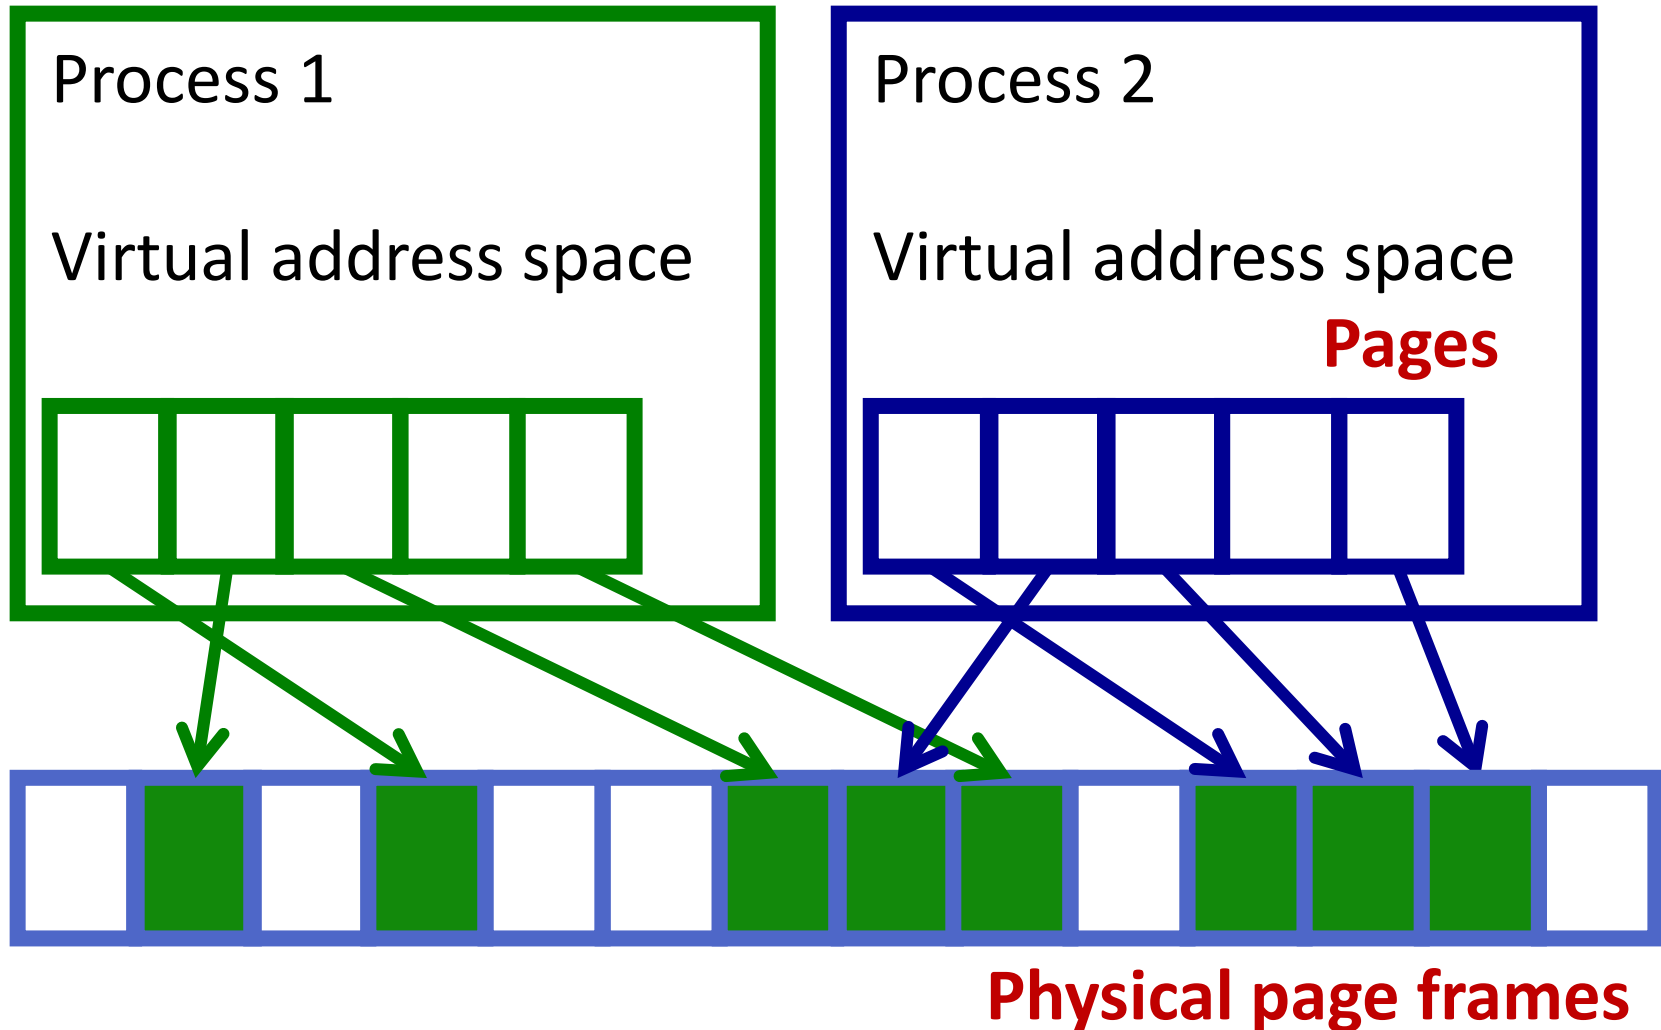
Virtual address space



# Paging



# Paging





# Advantages of paging

- **No** external fragmentation !
- Simplifies allocation, deallocation
- Provides application a contiguous view of memory
  - But, physical memory backing it could very scattered



# Advantages of paging

- **No** external fragmentation !
- Simplifies allocation, deallocation
- Provides application a contiguous view of memory
  - But, physical memory backing it could very scattered

Almost all processors today employs paging



# Disadvantages of paging

- **Internal fragmentation:** Memory wastage due to allocation only in page sizes (e.g., 4KB)
  - ▶ e.g., for an allocation 128 bytes a 4KB page will be allocated
  - ▶ Potentially waste (4KB – 128) bytes.

# Paging: How it works (simplistic)?

**Virtual address**

<b>0x0001bbf</b>	<b>21f</b>
------------------	------------

<b>Virtual Page Number (VPN)</b>	<b>Page Offset</b>
--	------------------------



# Paging: How it works (simplistic)?

**Virtual address**

<b>0x0001bbf</b>	<b>21f</b>
------------------	------------

**Virtual  
Page  
Number  
(VPN)**

**Page  
Offset**

<b>Prot</b>	<b>PFN</b>
--	-----
<b>R</b>	<b>0x30000</b>
<b>R/W</b>	<b>0x40100</b>
--	-----
<b>R/W</b>	<b>0x60000</b>

**Page table**

**In memory data structure  
maintained by the OS**



# Paging: How it works (simplistic)?

Virtual address

0x0001bbf	21f
-----------	-----

Virtual  
Page  
Number  
(VPN)

Page  
Offset

Prot

PFN

--	-----
R	0x30000
R/W	0x40100
--	-----
R/W	0x60000

Page table

In memory data structure  
maintained by the OS



# Paging: How it works (simplistic)?

Virtual address

0x0001bbf	21f
-----------	-----

Virtual  
Page  
Number  
(VPN)

Page  
Offset

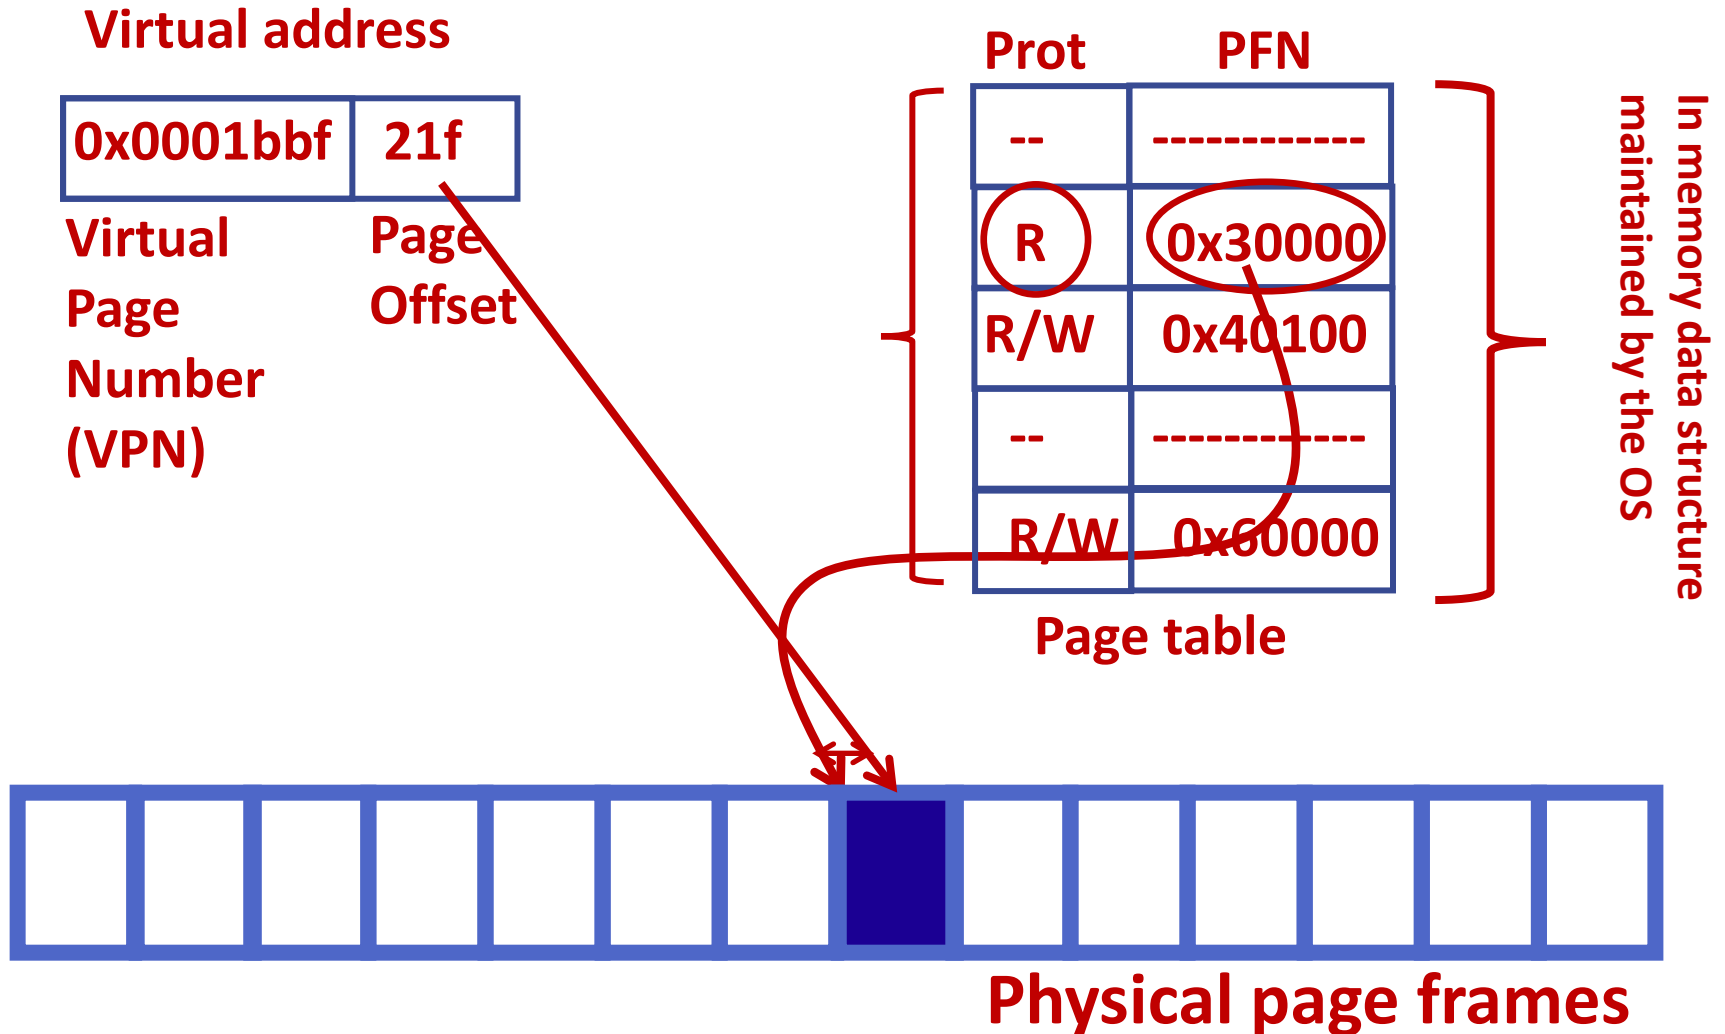
Prot	PFN
--	-----
<b>R</b>	<b>0x30000</b>
R/W	0x40100
--	-----
R/W	0x60000

Page table

In memory data structure  
maintained by the OS



# Paging: How it works (simplistic)?



# Paging: How it *really* works?

- Single level page table adds too much overhead
  - ▶ Consider a typical 48-bit virtual address space, 4KB page size and 8 byte long page table entry (PTE)
  - ▶ **Each** page table will be **512GB !**
  - ▶ There could be many processes → many page tables
  
- Often virtual address space is sparsely allocated
  - ▶ Entire address space is typically not allocated




# Paging: How it *really* works?

- Single level page table adds too much overhead
  - ▶ Consider a typical 48-bit virtual address space, 4KB page size and 8 byte long page table entry (PTE)
  - ▶ **Each** page table will be **512GB !**
  - ▶ There could be many processes → many page tables
- Often virtual address space is sparsely allocated
  - ▶ Entire address space is typically not allocated
- Solution: **Multilevel radix tree** for page table



# Paging: 64bit x86\* page table

4 level 512-ary radix tree indexed by virtual page number  
(latest x86-64 processors use 5-levels)



**Contains PFN**

\* i.e., Intel or AMD processors



# Paging: 64bit x86 page table

- Operating system maintains one page table per process (a.k.a., per virtual address space)
  - ▶ Allocated in normal cacheable memory – just like any other data structure
- Operating system creates, updates, deletes page table entries (PTEs)
  - ▶ Creates on demand/request and/or on process creation



# Paging: 64bit x86 page table

- Operating system maintains one page table per process (a.k.a., per virtual address space)
  - ▶ Allocated in normal cacheable memory – just like any other data structure
- Operating system creates, updates, deletes page table entries (PTEs)
  - ▶ Creates on demand/request and/or on process creation
- Page table structure is part of agreement between the OS and the hardware → page table structure is part of ISA (typically, e.g., x86 or ARM)




# Paging: Virtual to physical address translation

- Called *page table walk*.
  - ▶ Purpose: given a virtual address find the physical address?

# Paging: 64bit x86 page *walk*

**cr3**

Root of the page table.  
Part of process's context



Virtual page number

47

2 | 11


Page offset

0

# Paging: 64bit x86 page *walk*

**cr3**

Root of the page table.  
Part of process's context



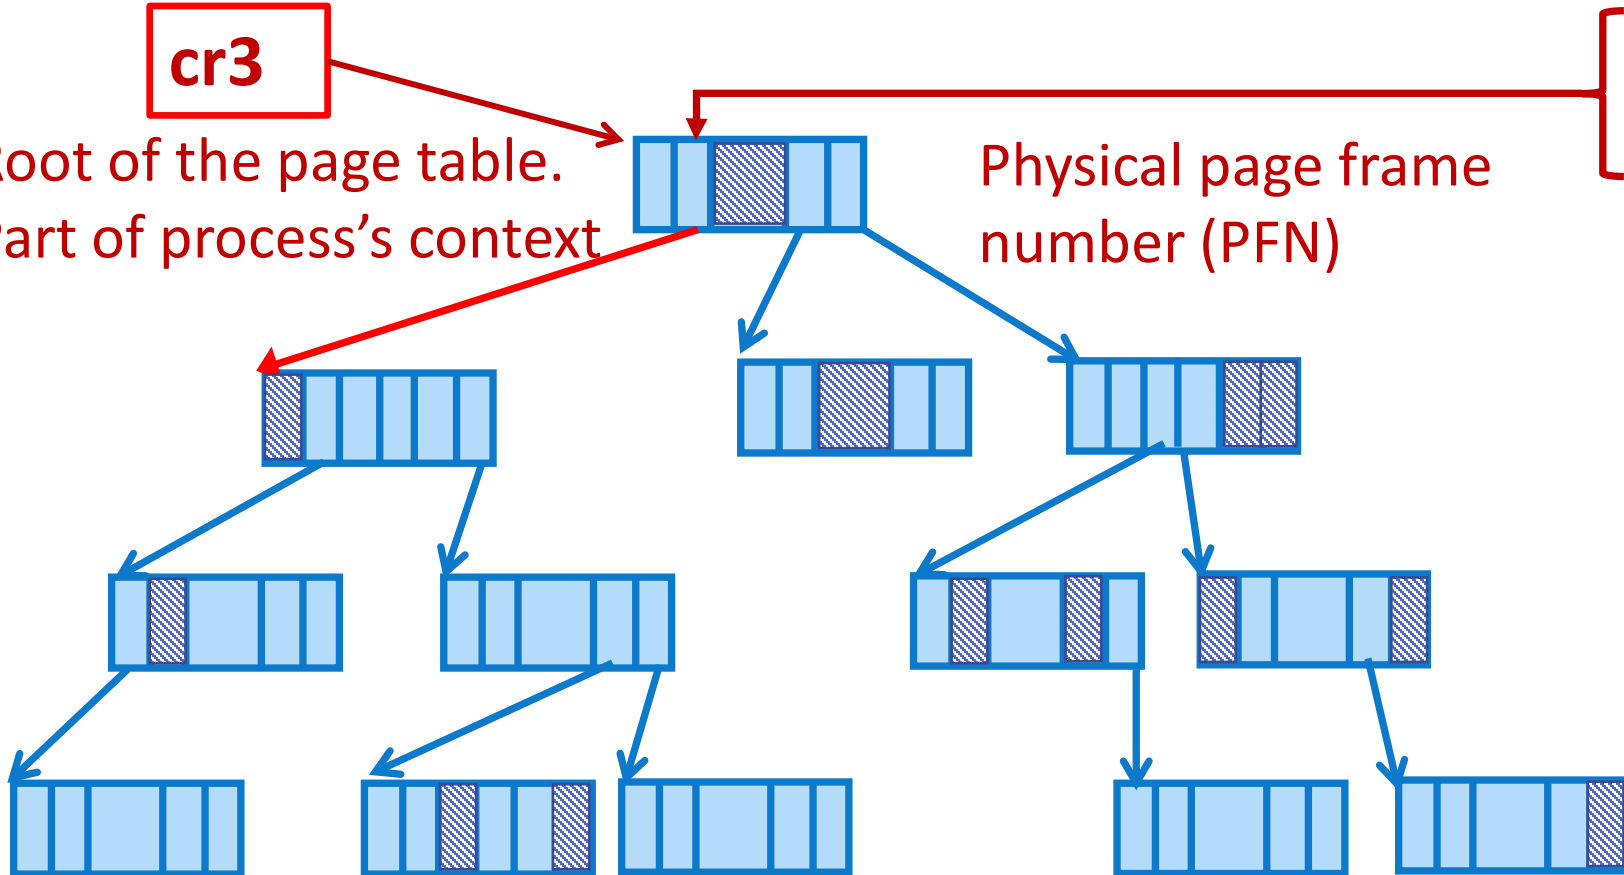
	47
	39
	38
	30
	29
	21
	20
	12
	11
Page offset	0

# Paging: 64bit x86 page *walk*

**cr3**

Root of the page table.  
Part of process's context

Physical page frame  
number (PFN)



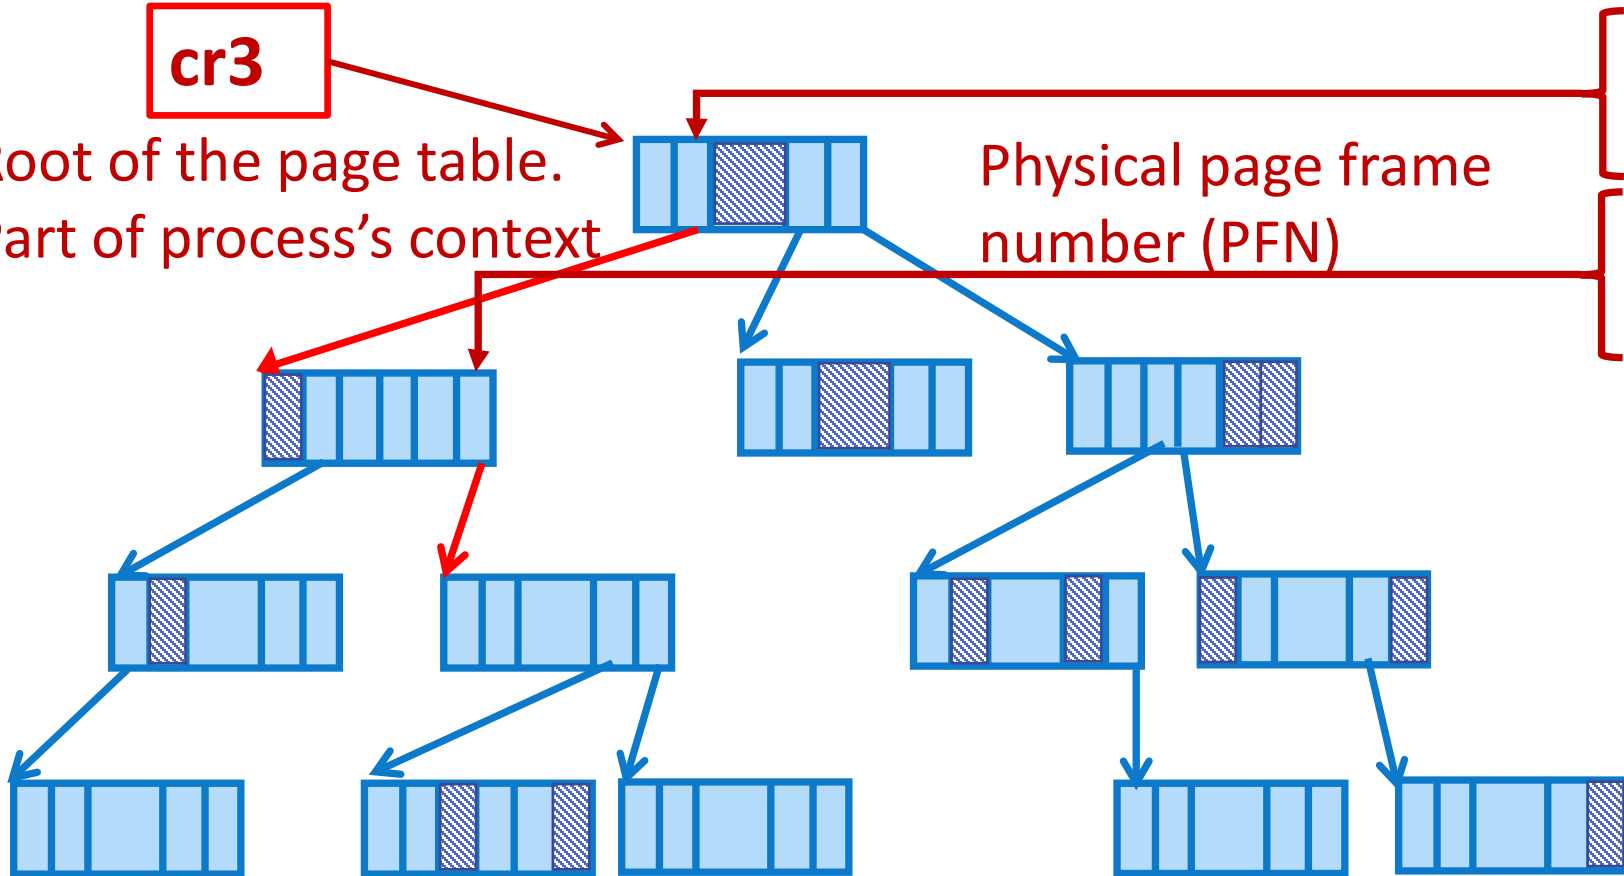
	47
	39
	38
	30
	29
	21
	20
	12
	11
Page offset	0

# Paging: 64bit x86 page *walk*

**cr3**

Root of the page table.  
Part of process's context

Physical page frame number (PFN)



	47
	39
	38
	30
	29
	21
	20
	12
	11
Page offset	0

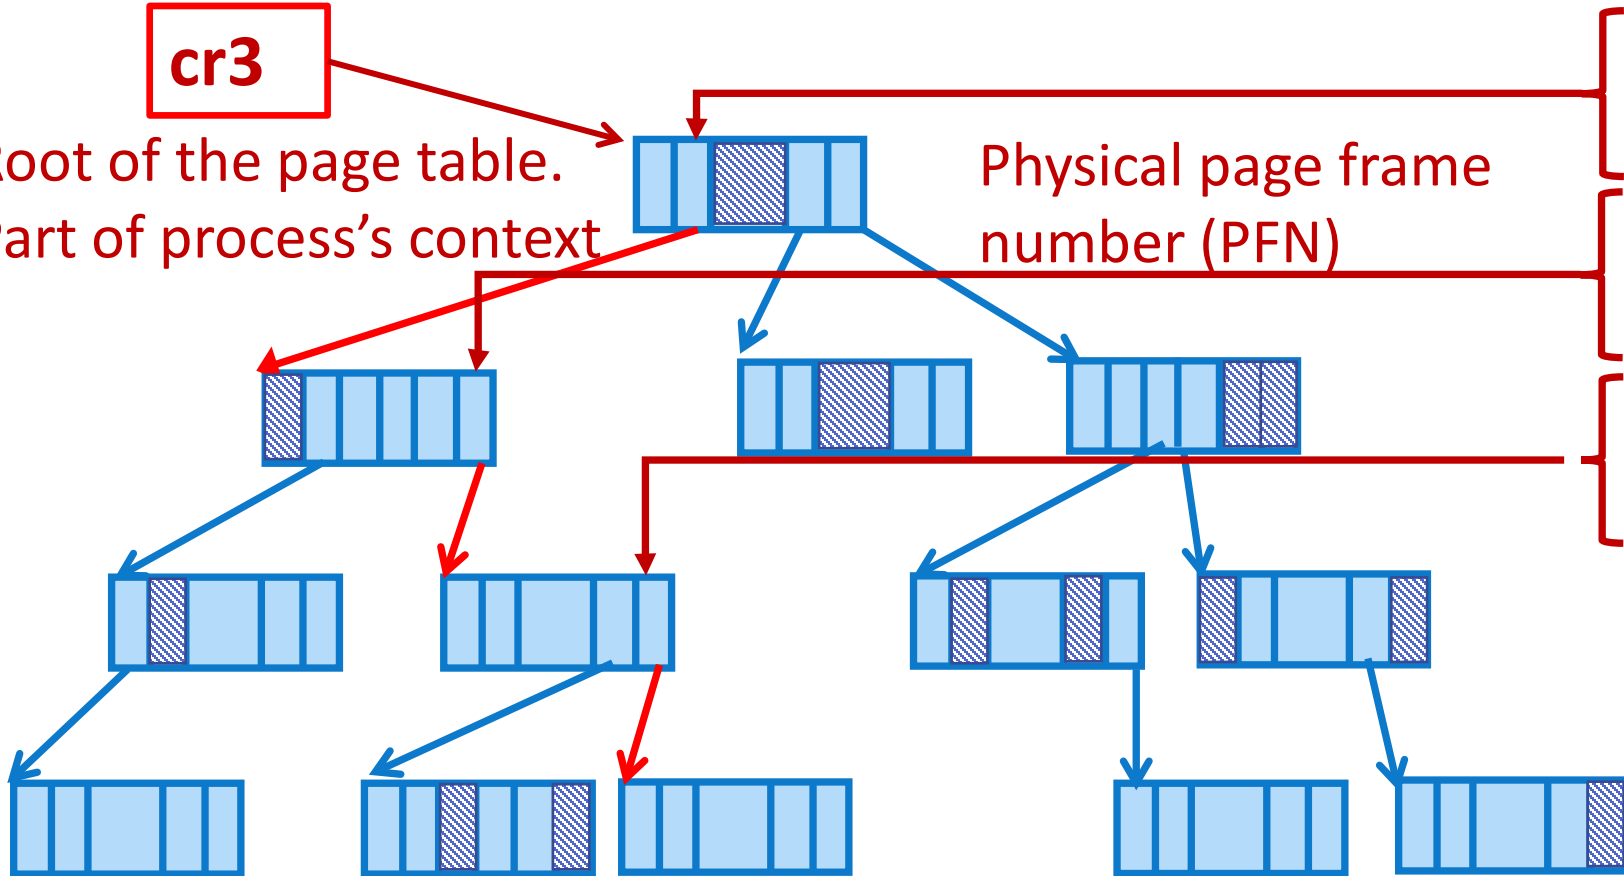


# Paging: 64bit x86 page *walk*

**cr3**

Root of the page table.  
Part of process's context

Physical page frame  
number (PFN)



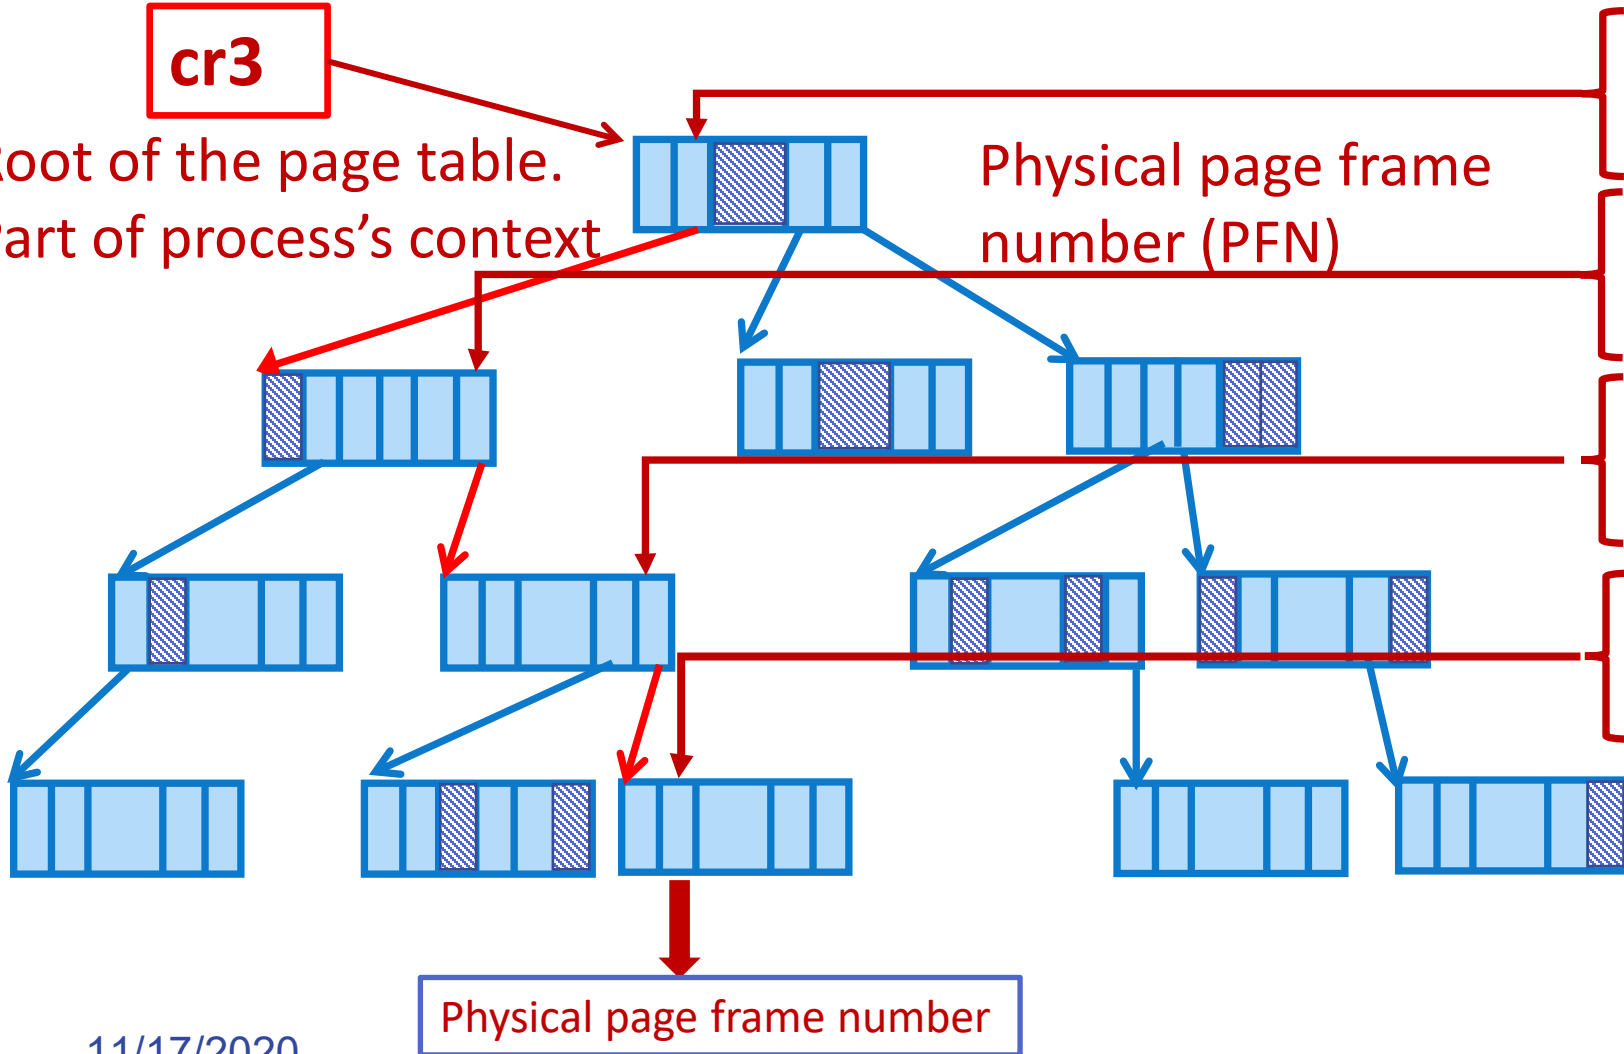
47	
39	
38	
30	
29	
21	
20	
12	
11	
0	Page offset

# Paging: 64bit x86 page *walk*

**cr3**

Root of the page table.  
Part of process's context

Physical page frame number (PFN)



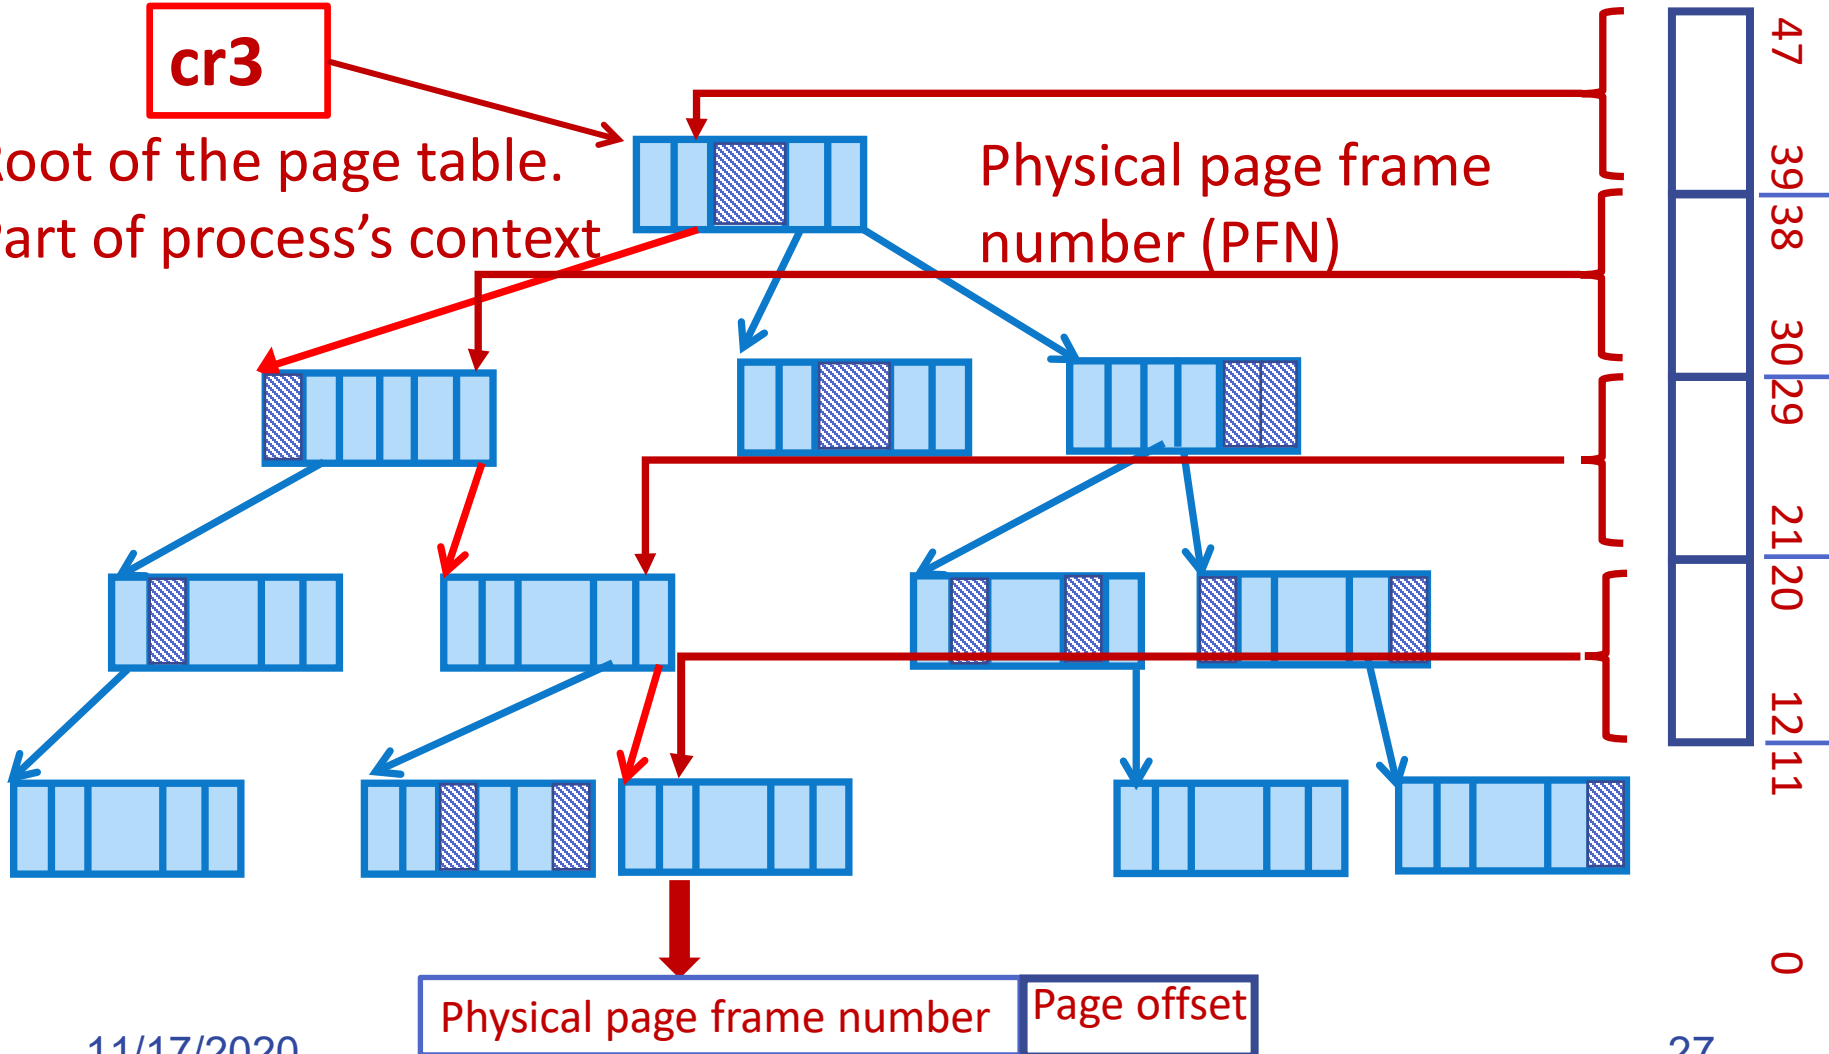
	47
	39
	38
	30
	29
	21
	20
	12
	11
Page offset	0

# Paging: 64bit x86 page *walk*

**cr3**

Root of the page table.  
Part of process's context

Physical page frame number (PFN)



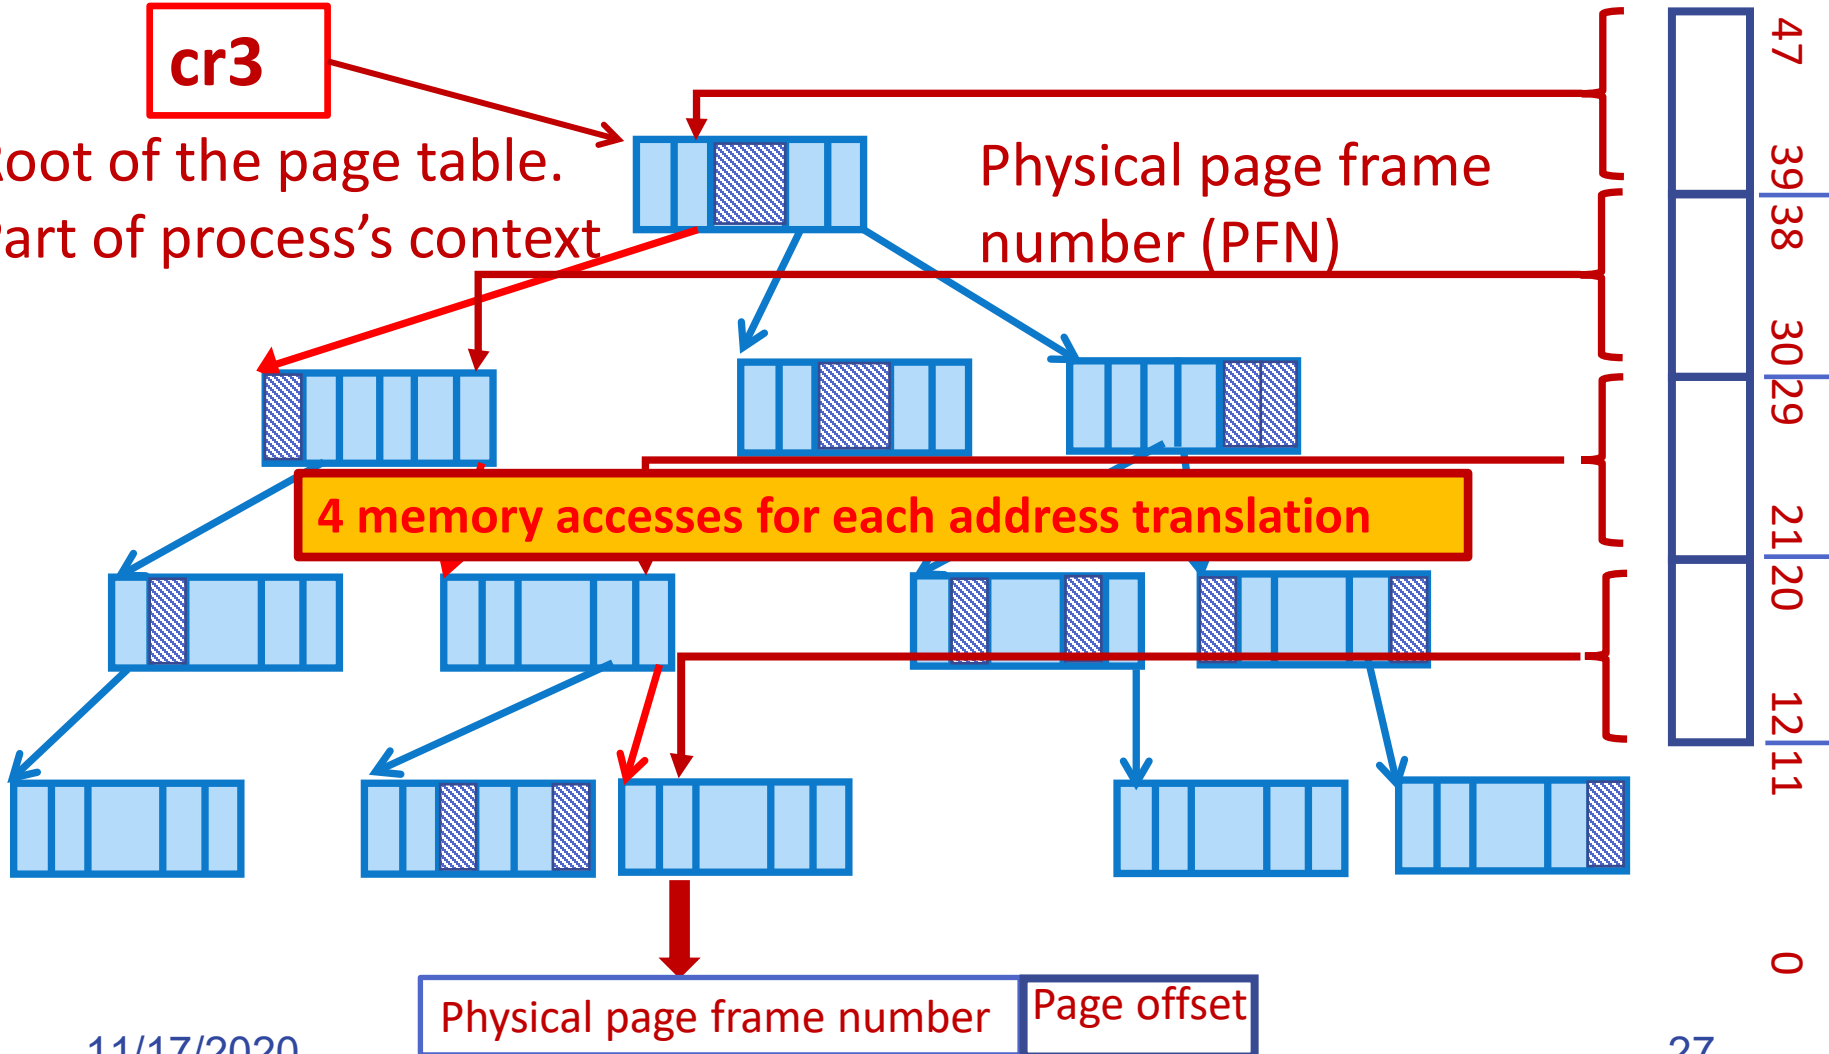
# Paging: 64bit x86 page *walk*

**cr3**

Root of the page table.  
Part of process's context

Physical page frame number (PFN)

**4 memory accesses for each address translation**





# Who walks the page table?

- A hardware page table walker (PTW) walks the page table (e.g., in Intel, AMD or ARM processors)
  - ▶ Input: Root of page table (cr3) and VPN
  - ▶ Output: Physical page frame number or a **page fault**
  - ▶ A hardware fine-state-automata in each CPU core
  - ▶ Generates **load-like “instructions”** to access addresses containing the memory holding desired page table entries

# Who walks the page table?

- A hardware page table walker (PTW) walks the page table (e.g., in Intel, AMD or ARM processors)
  - ▶ Input: Root of page table (cr3) and VPN
  - ▶ Output: Physical page frame number or a **page fault**
  - ▶ A hardware fine-state-automata in each CPU core
  - ▶ Generates **load-like “instructions”** to access addresses containing the memory holding desired page table entries
- Alternatives: Software page table walker
  - ▶ A OS handler walks the page table



# Who walks the page table?

- A hardware page table walker (PTW) walks the page table (e.g., in Intel, AMD or ARM processors)
  - ▶ Input: Root of page table (cr3) and VPN
  - ▶ Output: Physical page frame number or a **page fault**
  - ▶ A hardware fine-state-automata in each CPU core
  - ▶ Generates **load-like “instructions”** to access addresses containing the memory holding desired page table entries
- Alternatives: Software page table walker
  - ▶ A OS handler walks the page table
  - ▶ Advantage: Free to choose page table format
  - ▶ Disadvantage: Slow → Large address translation overhead
  - ▶ Example: SPARC (Sun/Oracle) machines



# TLB: Making page walks faster

- **Disadvantage:** A single address translation can take **4 memory accesses!**
  - ▶ To access one byte, 4 memory accesses needed for address translation



# TLB: Making page walks faster

- **Disadvantage:** A single address translation can take **4 memory accesses!**
  - ▶ To access one byte, 4 memory accesses needed for address translation
- How to make address translation fast?
  - ▶ **Translation Lookaside Buffer** or TLB to cache recently used virtual to physical address mappings



# TLB: Making page walks faster


- **Disadvantage:** A single address translation can take **4 memory accesses!**
  - ▶ To access one byte, 4 memory accesses needed for address translation
- How to make address translation fast?
  - ▶ **Translation Lookaside Buffer** or TLB to cache recently used virtual to physical address mappings
  - ▶ A read-only cache of contents of page table entries
  - ▶ For address translation of every load/store, TLB is first looked up
  - ▶ On a TLB **miss** page walk is performed
  - ▶ A TLB hit is fast but page walk is slow




# Typical TLB hierarchy of modern processors

- **Each core** typically has:
  - ▶ 32-64-entry L1 TLB (fully associative/4-8 way set-associative)
  - ▶ 1500-2500 entry L2 TLB (8-16 way set-associative)
  - ▶ Typically one page table walker (but latest Intel processor have two of them per core)
- **Note:** entire TLB hierarchy is private to a core (unlike cache subsystem)
  - ▶ Why?


# Where are TLBs & Page table walkers?




# Where are TLBs & Page table walkers?




# Where are TLBs & Page table walkers?



# Where are TLBs & Page table walkers?



# Contents of typical page table entry



Typical x86-64 PTE

- 'P' (Present bit): If the address present in memory
- 'U/S' (User/Supervisor bit): Is the page accessible in supervisor mode (e.g., by OS) only?
- 'R/W' (Read/Write): Is the page read-only?
- A (Access bit): Is this page has ever been accessed (load/stored to)?
- D (Dirty bit): Is the page has been written to?
- X/D (Executable bit): Does the page contains executable?



# Putting it all Together

