



# Memory consistency models and synchronizations



# Acknowledgements

- Several of the slides in the deck are from Luis Ceze (Washington), Nima Horanmand (Stony Brook), Mark Hill, David Wood, Karu Sankaralingam (Wisconsin), Abhishek Bhattacharjee (Rutgers).
- Development of this course is partially supported by Western Digital corporations.



# Coherence vs. Consistency

A=0    FLAG=0

C\_0

```
ST A 1;
ST FLAG 1;
```

C\_1

```
L1: LD r1 FLAG
If r1 == 0; JMP L1; // spin lock
LD r2 A;
```

# Coherence vs. Consistency

	A=0	FLAG=0	
<u>C_0</u>			<u>C_1</u>
ST A 1;			L1: LD r1 FLAG
ST FLAG 1;			If r1 == 0; JMP L1; // spin lock
			LD r2 A;

What value should r2 contain?

# Coherence vs. Consistency

	A=0	FLAG=0	
<u>C_0</u>			<u>C_1</u>
ST A 1;			L1: LD r1 FLAG
ST FLAG 1;			If r1 == 0; JMP L1; // spin lock
			LD r2 A;

What value should r2 contain? (expect r2=1)

But what does the coherence says?

# Coherence vs. Consistency

	A=0	FLAG=0	
<u>C_0</u>			<u>C_1</u>
ST A 1;			L1: LD r1 FLAG
ST FLAG 1;			If r1 == 0; JMP L1; // spin lock
			LD r2 A;

What value should r2 contain? (expect r2=1)

But what does the coherence says? Nothing

Coherence is about total ordering of stores and loads to one given memory location (in practice a single cache block)

Says nothing about ordering across different memory locations/addresses

# Coherence vs. Consistency

	A=0	FLAG=0	
<u>C_0</u>			<u>C_1</u>
ST A 1;			L1: LD r1 FLAG
ST FLAG 1;			If r1 == 0; JMP L1; // spin lock
			LD r2 A;

What value should r2 contain? (expect r2=1)

But what does the coherence says? Nothing

Coherence is about total ordering of stores and loads to one given memory location (in practice a single cache block)

Says nothing about ordering across different memory locations/addresses

But correctly specifying/writing multi-threaded program requires guarantees for ordering of load/stores to different locations



# Memory consistency models

A memory consistency model specifies global orders of writes to **all memory** locations relative to each other

- A memory consistency model tells what are the legal reordering of loads/stores to different memory locations
- Different memory consistency models possible
  - Tradeoff between ease of programmability vs. performance
  - "Relaxed" models enforces less ordering (better performance) but harder to program
- Memory consistency model part of ISA
  - X86 and ARM has different memory consistency model





# Memory consistency models

- Specifies which re-ordering of loads and stores (to different addresses) are allowed

	A=0	FLAG=0	
<u>C_0</u>			<u>C_1</u>
ST A 1;			L1: LD r1 FLAG
ST FLAG 1;			If r1 == 0; JMP L1; // spin lock
			LD r2 A;



# Memory consistency models

- Specifies which re-ordering of loads and stores (to different addresses) are allowed

	A=0	FLAG=0	
<u>C_0</u>			<u>C_1</u>
ST A 1;			L1: LD r1 FLAG
ST FLAG 1;			If r1 == 0; JMP L1; // spin lock
			LD r2 A;

- For example, this program will work fine if load, store bypassing is not allowed within a given thread, even if they are to different addresses



# Memory consistency models

- Specifies which re-ordering of loads and stores (to different addresses) are allowed

```

                                A=0   FLAG=0
C_0                                C_1
ST A 1;                          L1: LD r1 FLAG
ST FLAG 1;                       If  r1 == 0; JMP L1; // spin lock
                                LD r2 A;
```

- For example, this program will work fine if load, store bypassing is not allowed within a given thread, even if they are to different addresses
- Parallel programmers rely on the memory model to reason about correctness of your program



# Who defines memory consistency models?


- Any programmer of parallel shared memory programs should care of memory consistency models
- Defined by the H/W –described in ISA
- Compilers should also need to define their memory consistency models
  - ▶ Determines what optimizations are allowed



# Many possible memory consistency models

- Sequential consistency
  - ▶ Most intuitive, programmer friendly
  - ▶ But most restriction in terms of performance optimizations
  - ▶ e.g., implemented in MIPS R10K
- Total store order (processor consistency model)
  - ▶ Less restrictive than
  - ▶ Intel, AMD's x86'64 processors
  - ▶ Sun UltraSPARC
- Weak memory models (release consistency models)
  - ▶ Most relaxed, burdens programmer but allows more hardware optimizations
  - ▶ ARM, IBM POWER processors

# Sequential consistency




**Per-processor program order:** memory operations from individual processors maintain program order

**Single sequential order:** the memory operations from all processors maintain a single sequential order

[Lamport'79]


# Sequential consistency



**Per-processor program order:** memory operations from individual processors maintain program order

**Single sequential order (Total order):** memory operations from all processors maintain a single sequential order

# Sequential consistency



C1

st A  
ld C  
st C

C2

st A  
st C  
ld D

A possible legal  
global order


st A  
st A  
ld C  
st C  
st C  
ld D

**Per-processor program order:** memory operations from individual processors maintain program order

**Single sequential order (Total order):** memory operations from all processors maintain a single sequential order



# Sequential consistency



C1

st A  
ld C  
st C

C2

st A  
st C  
ld D

A possible legal global order

st A	st A
st A	ld C
ld C	st A
st C	st C
st C	st C
ld D	ld D

**Per-processor program order:** memory operations from individual processors maintain program order

**Single sequential order (Total order):** memory operations from all processors maintain a single sequential order



# Sequential consistency (SC)

	A=0	FLAG=0	
<u>C_0</u>			<u>C_1</u>
ST A 1;			L1: LD r1 FLAG
ST FLAG 1;			If r1 == 0; JMP L1; // spin lock
			LD r2 A;

- What will be the value loaded in r2 under SC?

# Reordering of load/stores

Program order	Earlier operation	LD	ST
	Later operation		
	LD	NO	NO
	ST	NO	NO

Allowed **reordering** of load/stores to **different addresses** under sequential consistency.

No reordering for load/stores to same address – ensured by coherence

# Food for thought (assume SC)

- Answer the following questions:
  - Initially: all variables zero (that is, x is 0, y is 0, flag is 0, A is 0)
  - What value pairs can be read by the two loads? (x, y) pairs:

Core 0

Core 1

LD x	ST y 1
LD y	ST x 1